

Fast-DetectGPT

Understanding sampling-based zero-shot
AI-generated Text Detection

Krystof Bobek
Mara Dragomir
Mahdi Rahimi
Jose Garcia
Salvador Torpes

Table of Contents

1. Introduction
2. Methodology
 1. Models
 2. Dataset Generation
 3. Methods
 4. Sampling and Scoring
3. Experiments and Results
 1. New Datasets & Models
 2. Mid-K Adversarial Attack
4. Conclusion

Introduction

The Task

Classifying input text as human or AI-generated

With roots in blues rock and psychedelic rock, the bands that created heavy metal developed a thick, massive sound, characterized by highly amplified distortion, extended guitar solos, emphatic beats, and overall loudness.

Which example is human and which is generated by a LLM?

Metal music is a genre of rock music that typically features heavy, distorted guitar riffs and fast-paced tempos. It often incorporates elements of other musical genres such as punk, hardcore, and classical music, and is known for its aggressive, energetic sound.

The Task

Classifying input text as human or AI-generated

Human

With roots in blues rock and psychedelic rock, the bands that created heavy metal developed a thick, massive sound, characterized by highly amplified distortion, extended guitar solos, emphatic beats, and overall loudness.

LLM

Metal music is a genre of rock music that typically features heavy, distorted guitar riffs and fast-paced tempos. It often incorporates elements of other musical genres such as punk, hardcore, and classical music, and is known for its aggressive, energetic sound.

Difficult task for humans, but maybe not so much for AI


Motivation

Why is it important to detect AI-generated content?

- Preventing mis- and disinformation
- Preventing plagiarism
- Transparency
- Preserving human creativity and intellect

The Idea

DetectGPT & Fast DetectGPT



The probability curvature of a text sample fed to a model can show whether the text is written by a human or by the model itself

Methodology

Models

Models

FastDetectGPT uses three different Large Language Models in its pipeline:

Source Model

Source model is the one used to generate the datasets: we feed the model with the first 30 tokens of a human written sentence and then we do decoding up until we have a sentence that is the size of the original one.

Sampling Model

Sampling model is the one used to get the probability distribution of tokens over each position in the sentence that will then be used to sample from and generate the 'perturbations' of the sentence

Scoring Model

Scoring model is the one used to get the probability distribution of tokens over each position in the sentence that will then be used to compute the loss of the sentence

Models

FastDetectGPT uses three different Large Language Models in its pipeline:

Source Model

- Used for dataset creation
- Takes a collection of human sentences as input, and generates a corresponding AI sentence for each using the first 30 tokens of the sentence

Sampling Model

- Used for perturbing the input sentences
- Samples from the probability distribution of tokens over each position in the sentence

Scoring Model

- Used for the probability curvature
- Uses the probability distribution of tokens over each position in the sentence to compute the loss
- Classifies the sentence via a threshold

Settings

The experiments of FastDetectGPT can be done under a black box or a white box setting.

Black Box

An experiment under the black box setting is defined by having **different** source and scoring models - real life setting where we don't know if the text was AI generated

White Box

An experiment under the white box setting is defined by having the **same** source and scoring models - we know the model we are trying to detect

In both settings sampling and scoring model may differ and that might lead to performance changes.

Settings

The experiments of FastDetectGPT can be done under a black box or a white box setting.

(Our Focus)

Black Box

where the source and scoring models can be different (the model doesn't know what model generated the text)

White Box

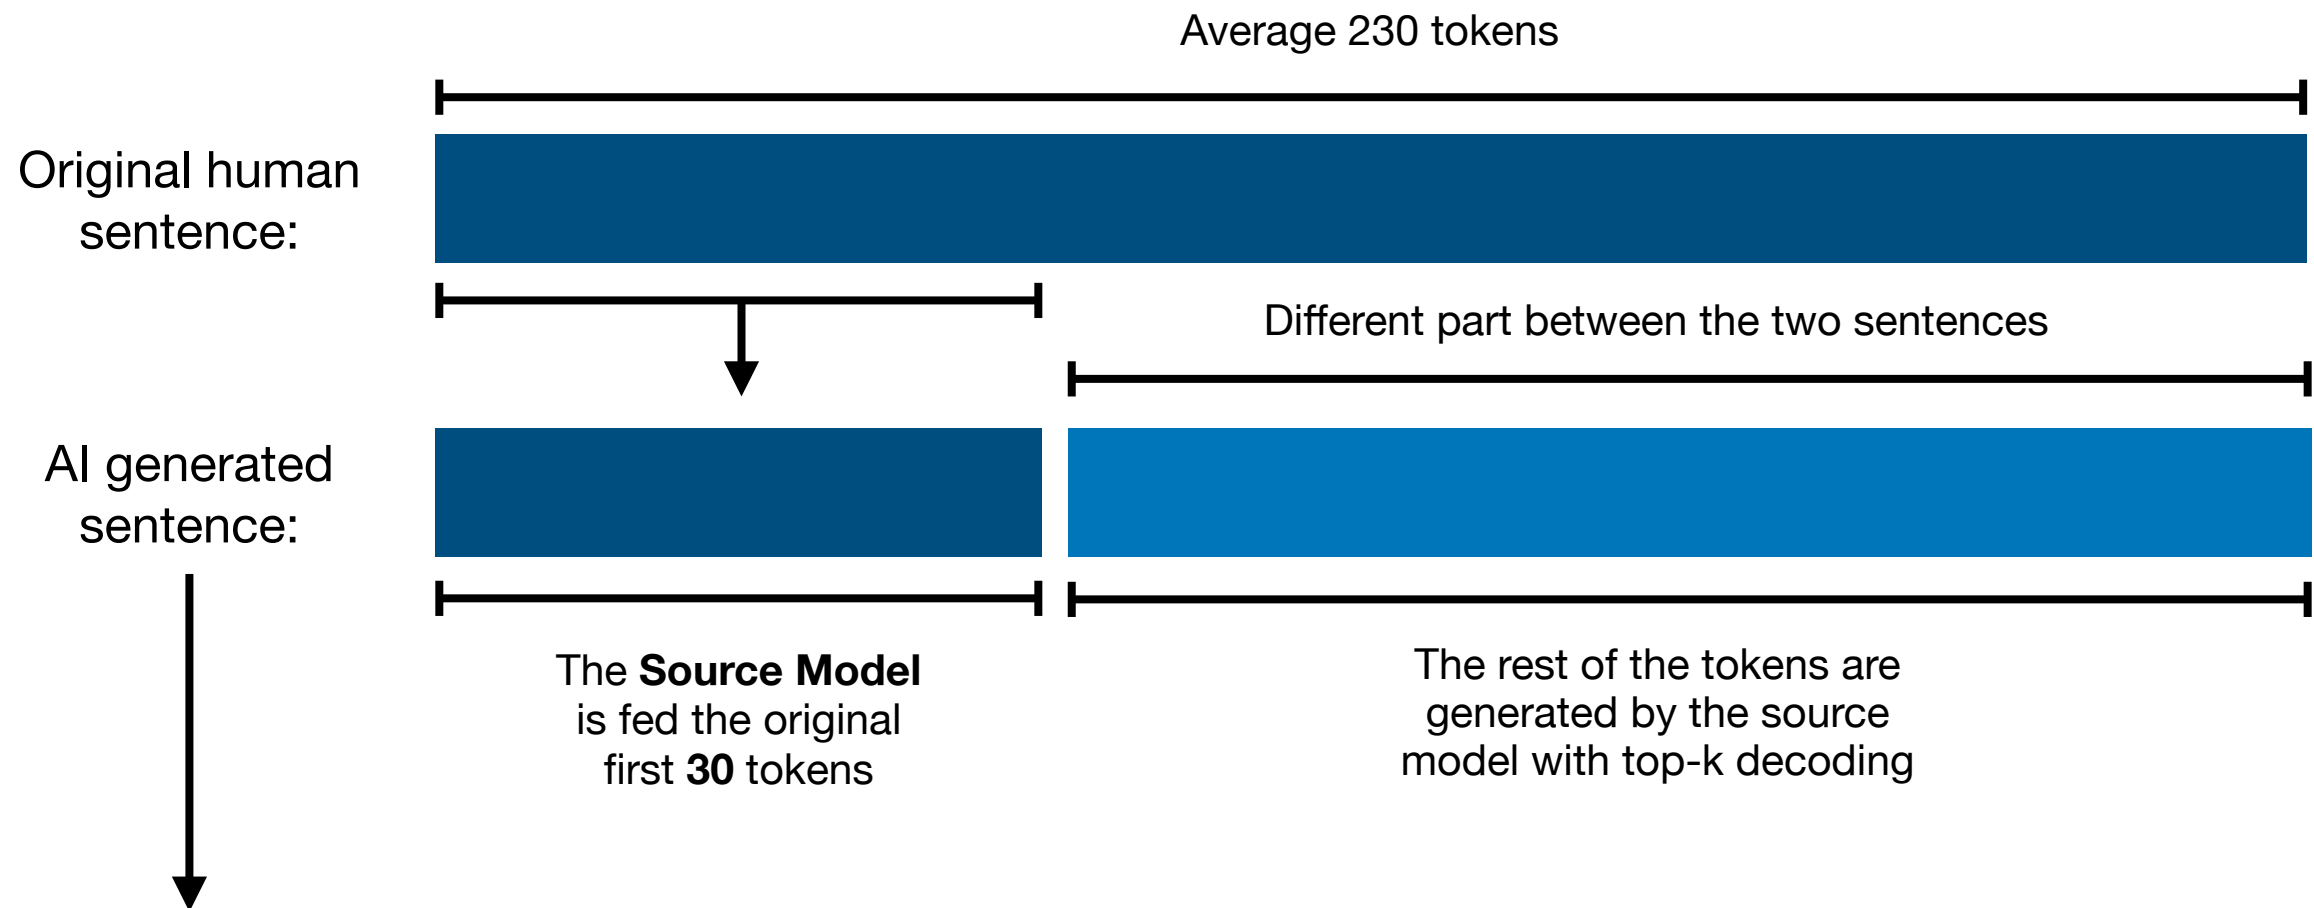
where the source and scoring models are the same (the model classifies whether a text is generated by itself)

The sampling and scoring model may differ and that might lead to performance changes

Dataset Generation

Generating AI text

Considering we have a human text dataset with multiple entries, building the AI text will consist of building an AI equivalent sentence for each entry, as follows:

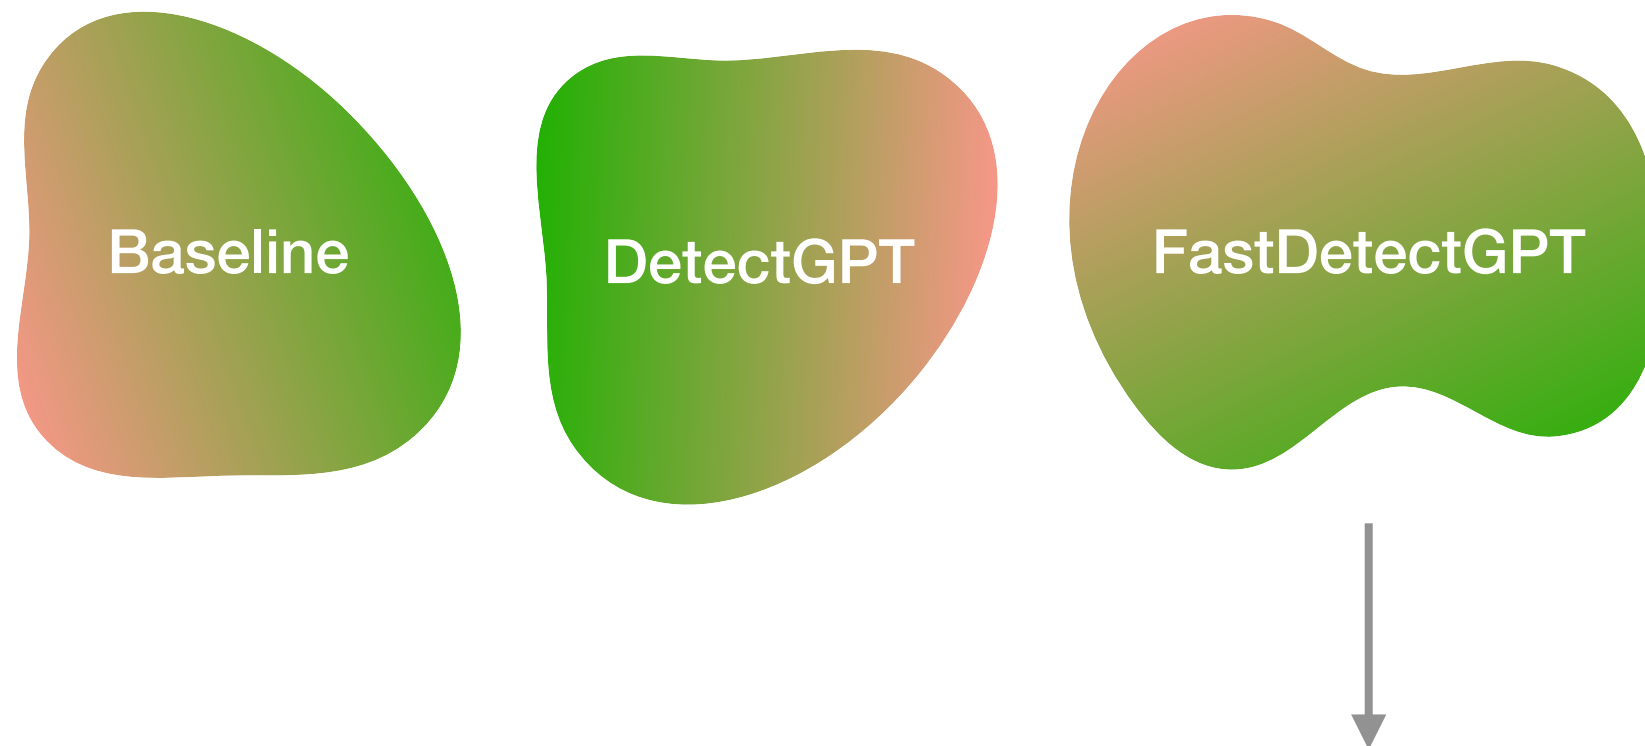


The AI generated sentence will be fed into one of the three models to conclude if it's AI generated or not: baseline models, DetectGPT or FastDetectGPT

Methods

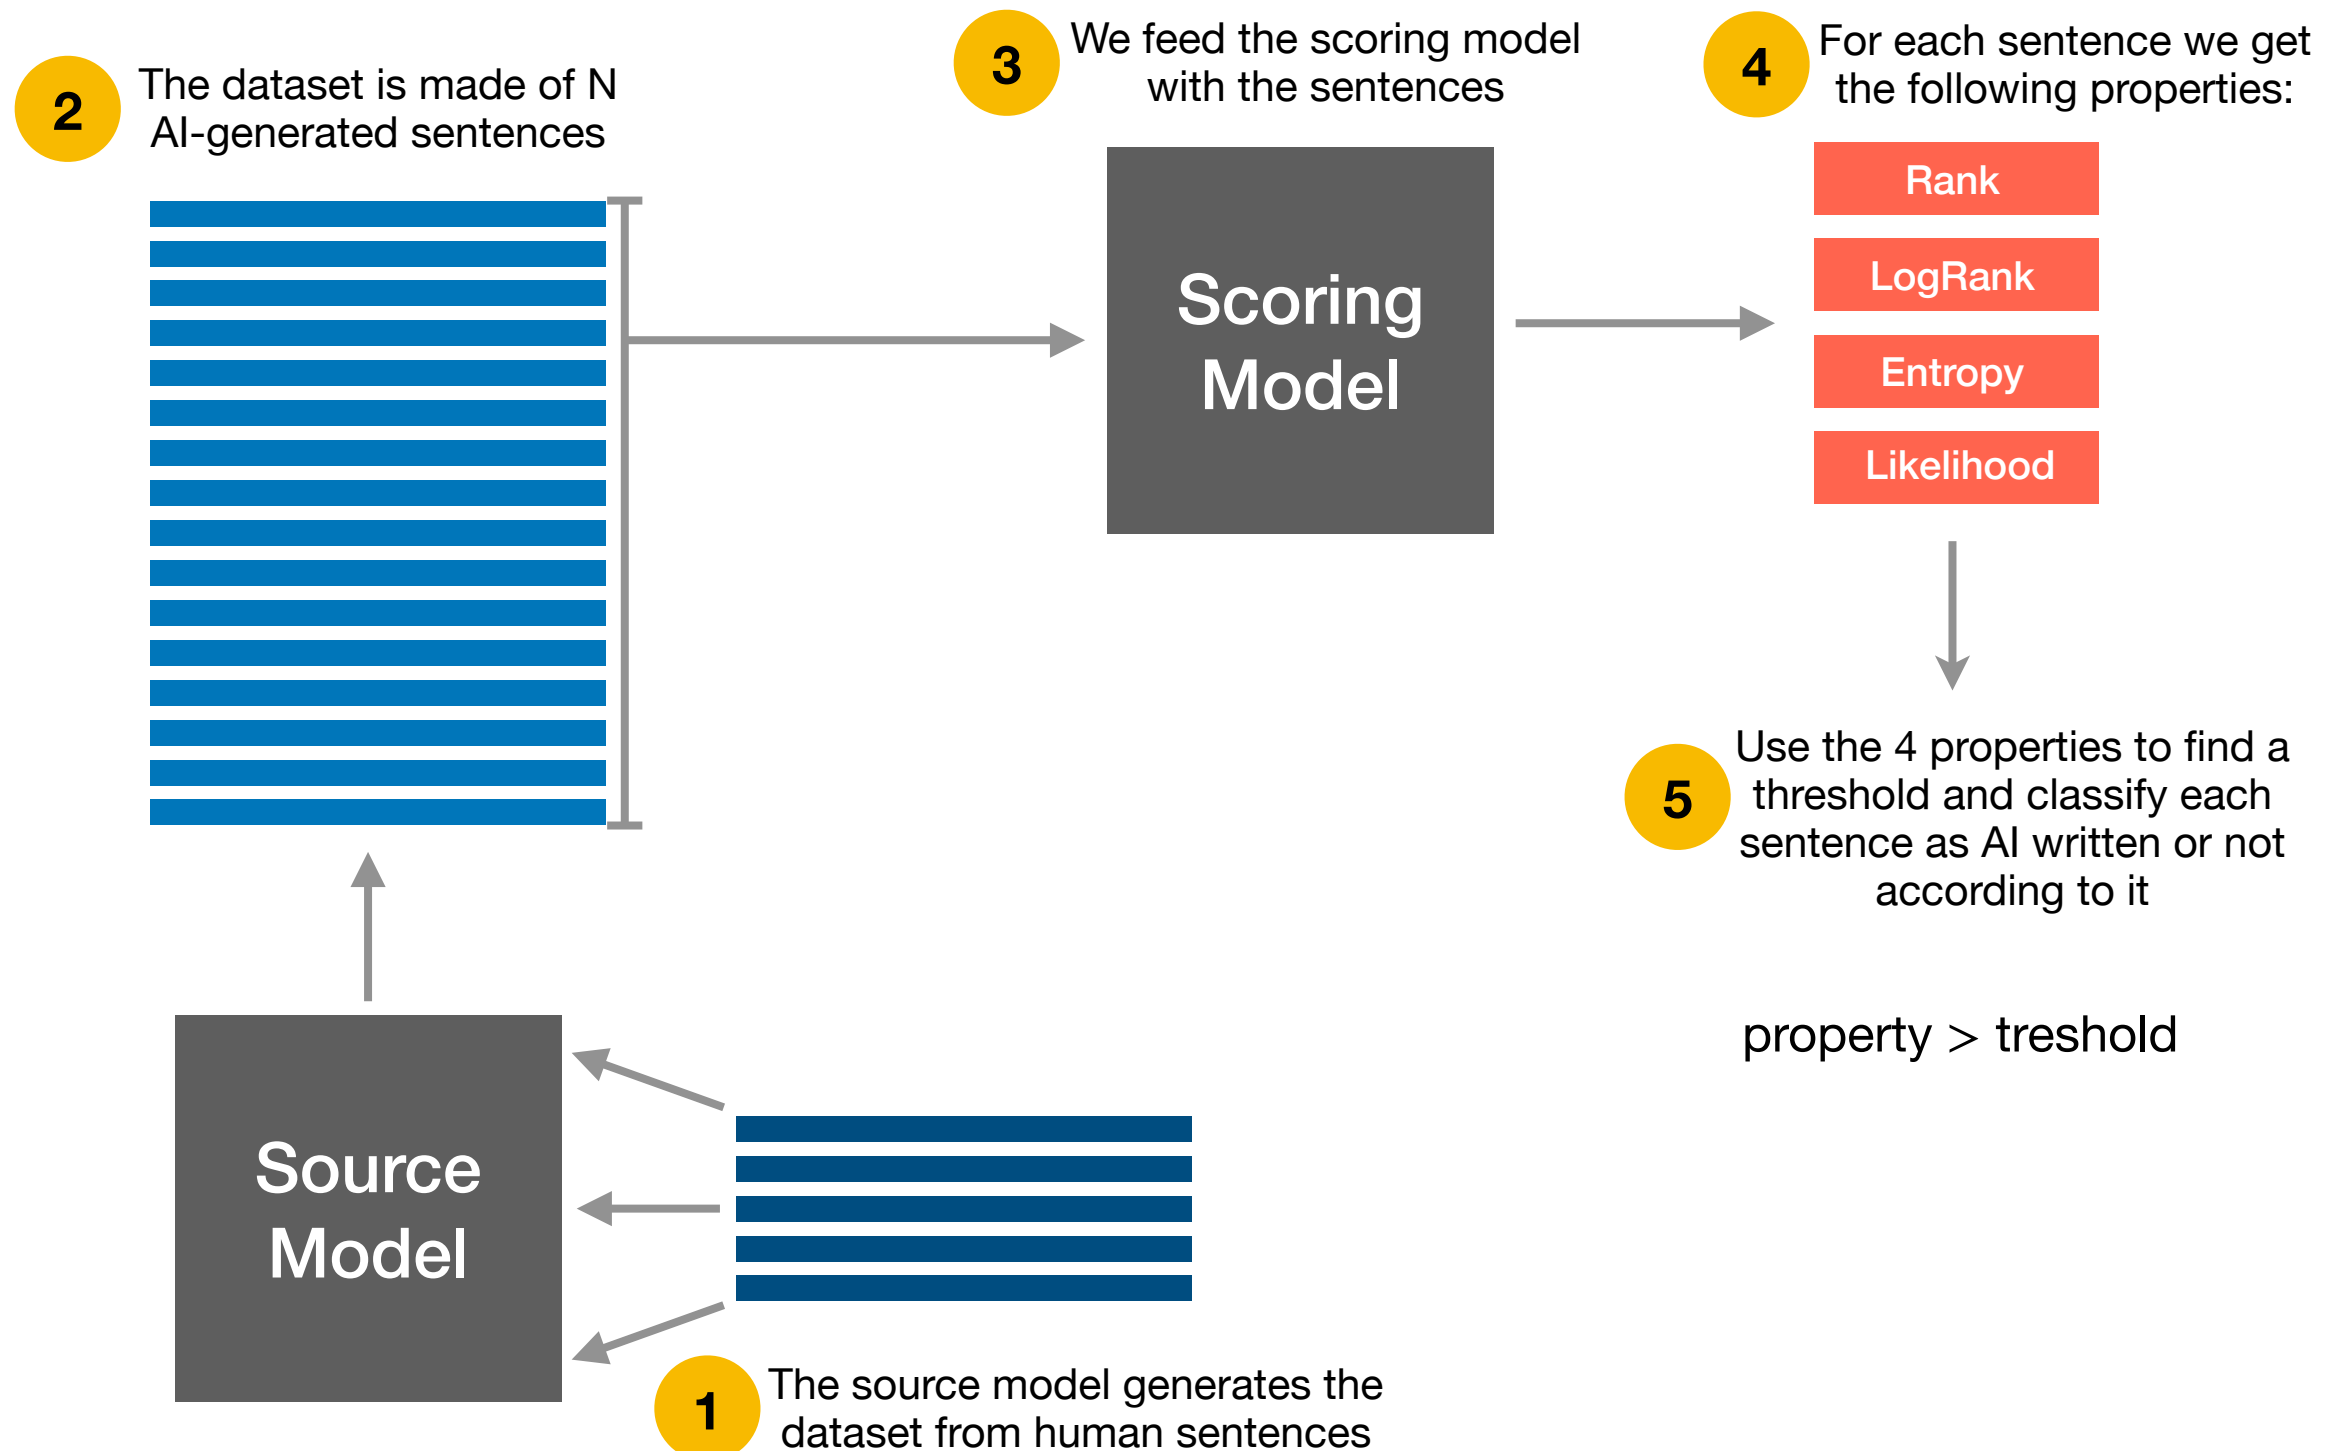
Three Different Methods for Detection

We have studied three different methods for identifying AI-written text:

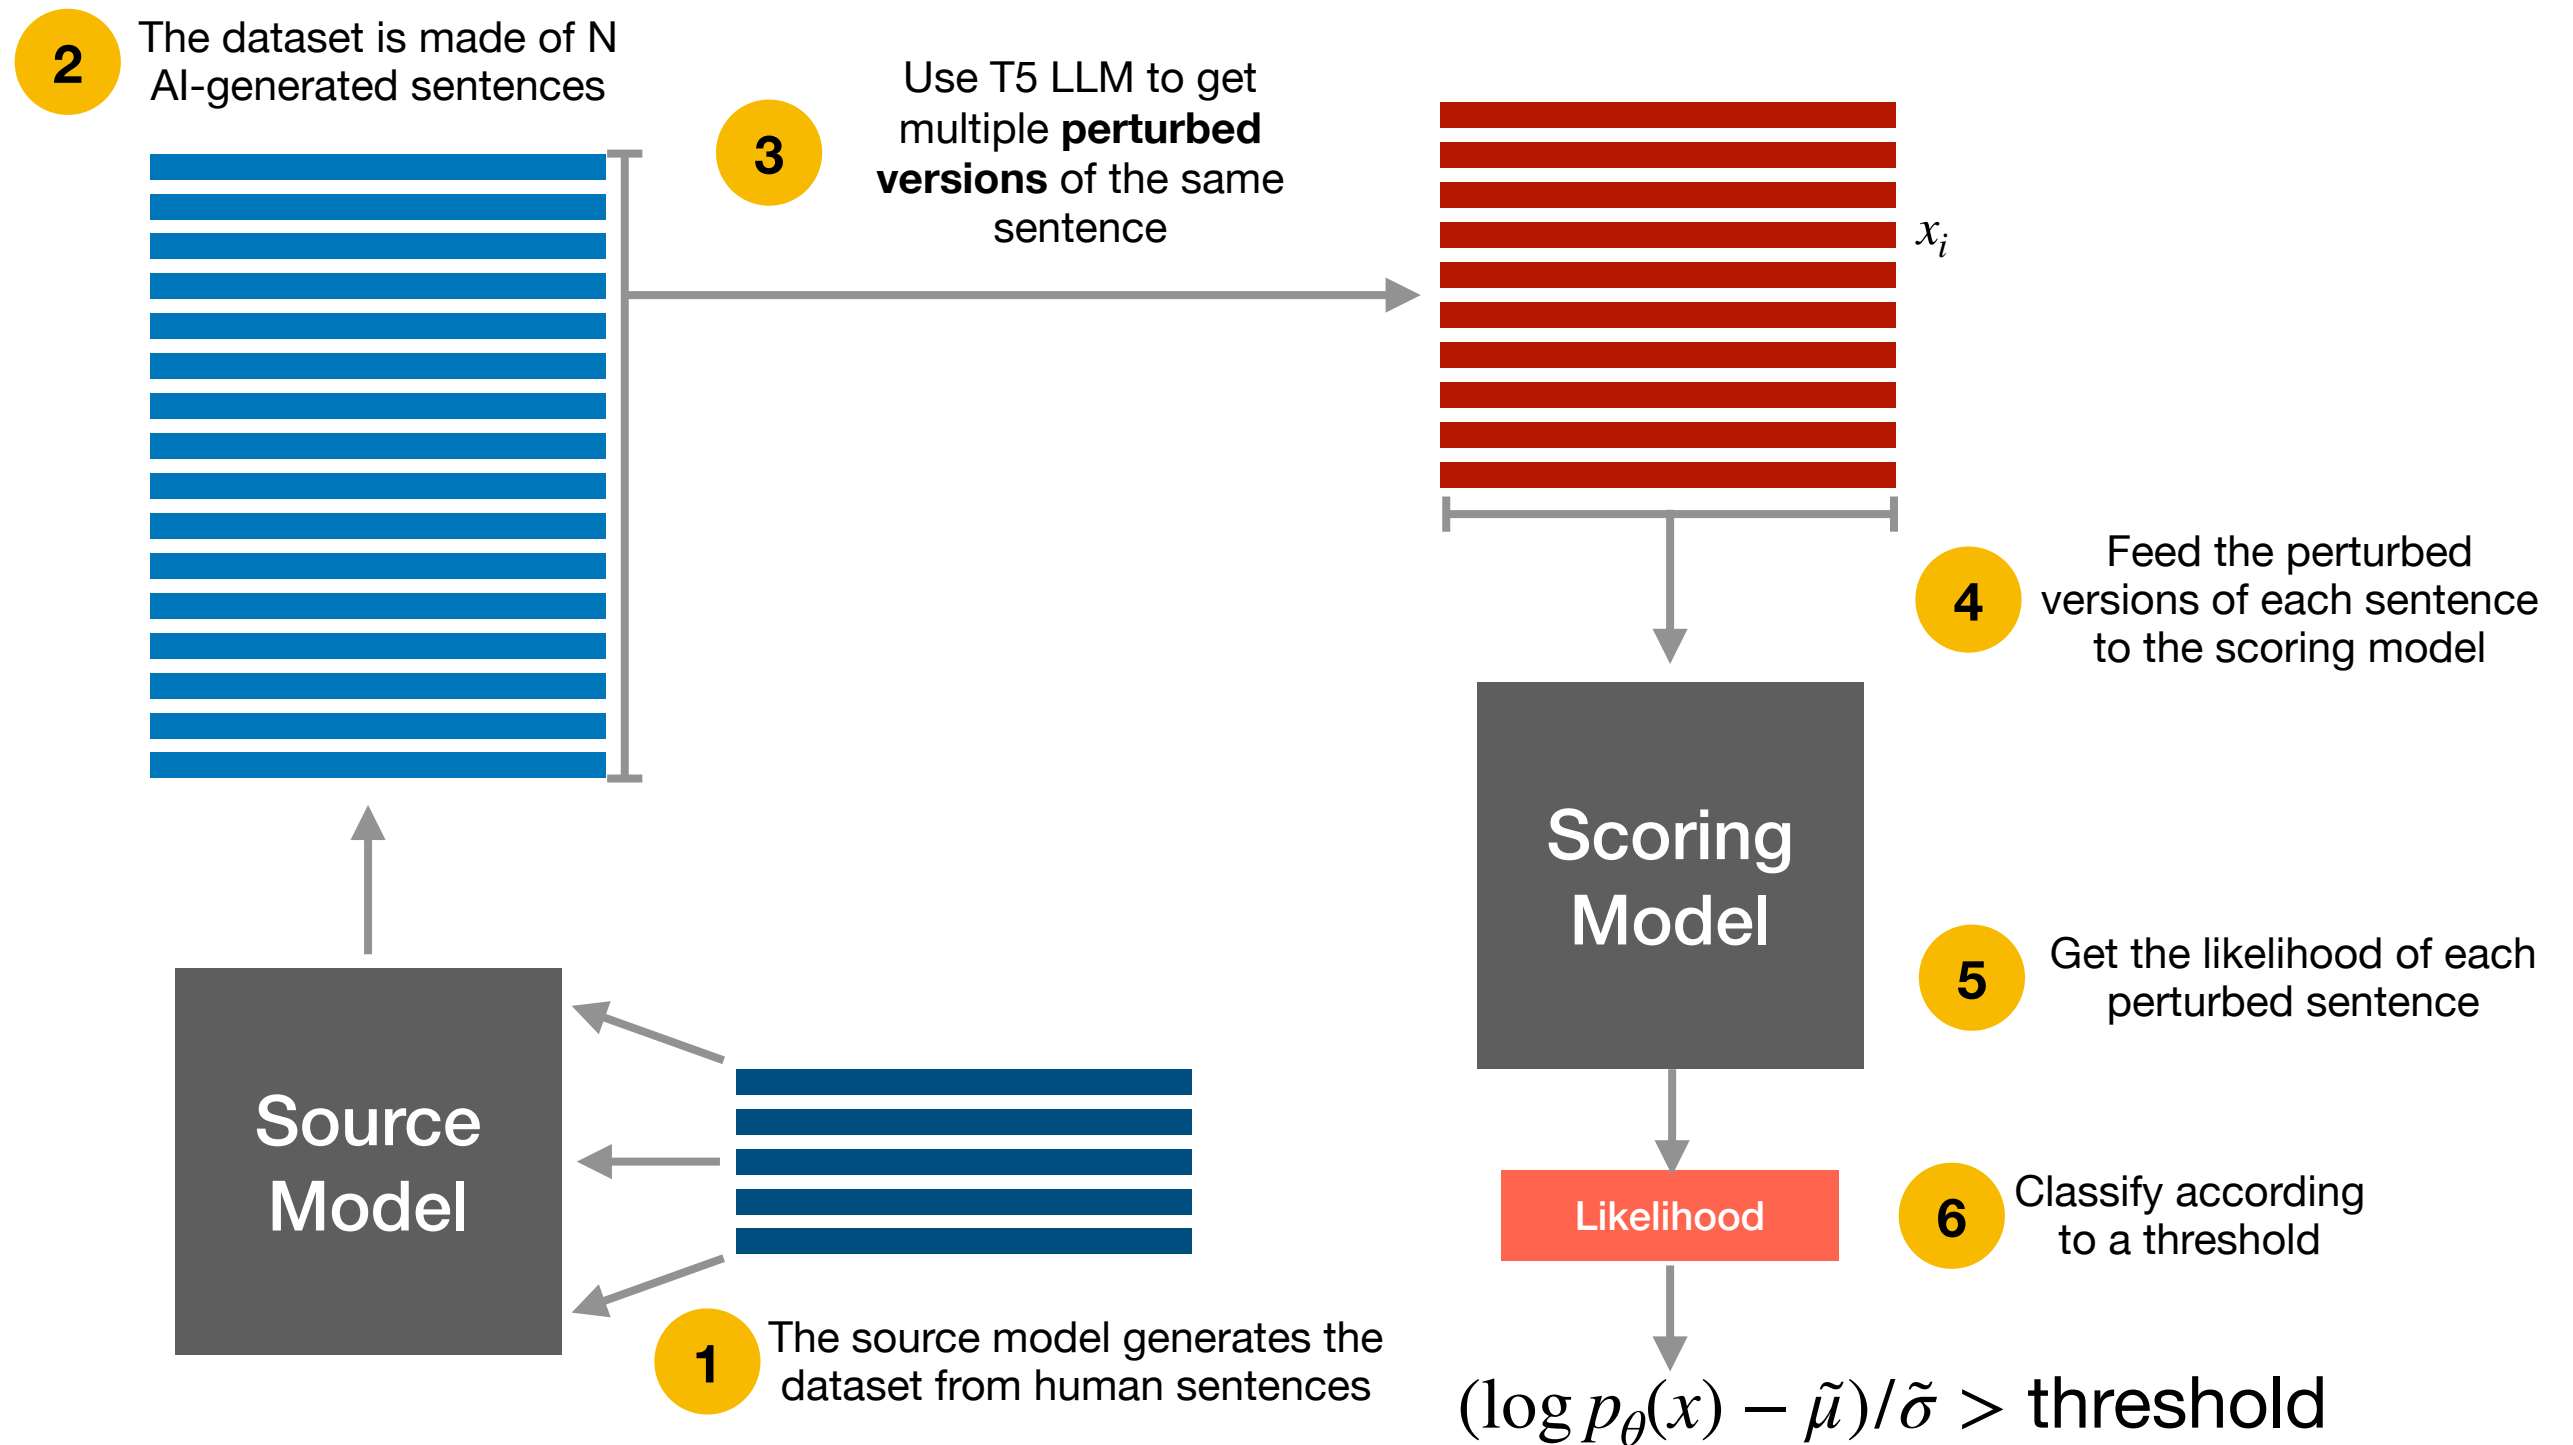


Our work is focused on this one

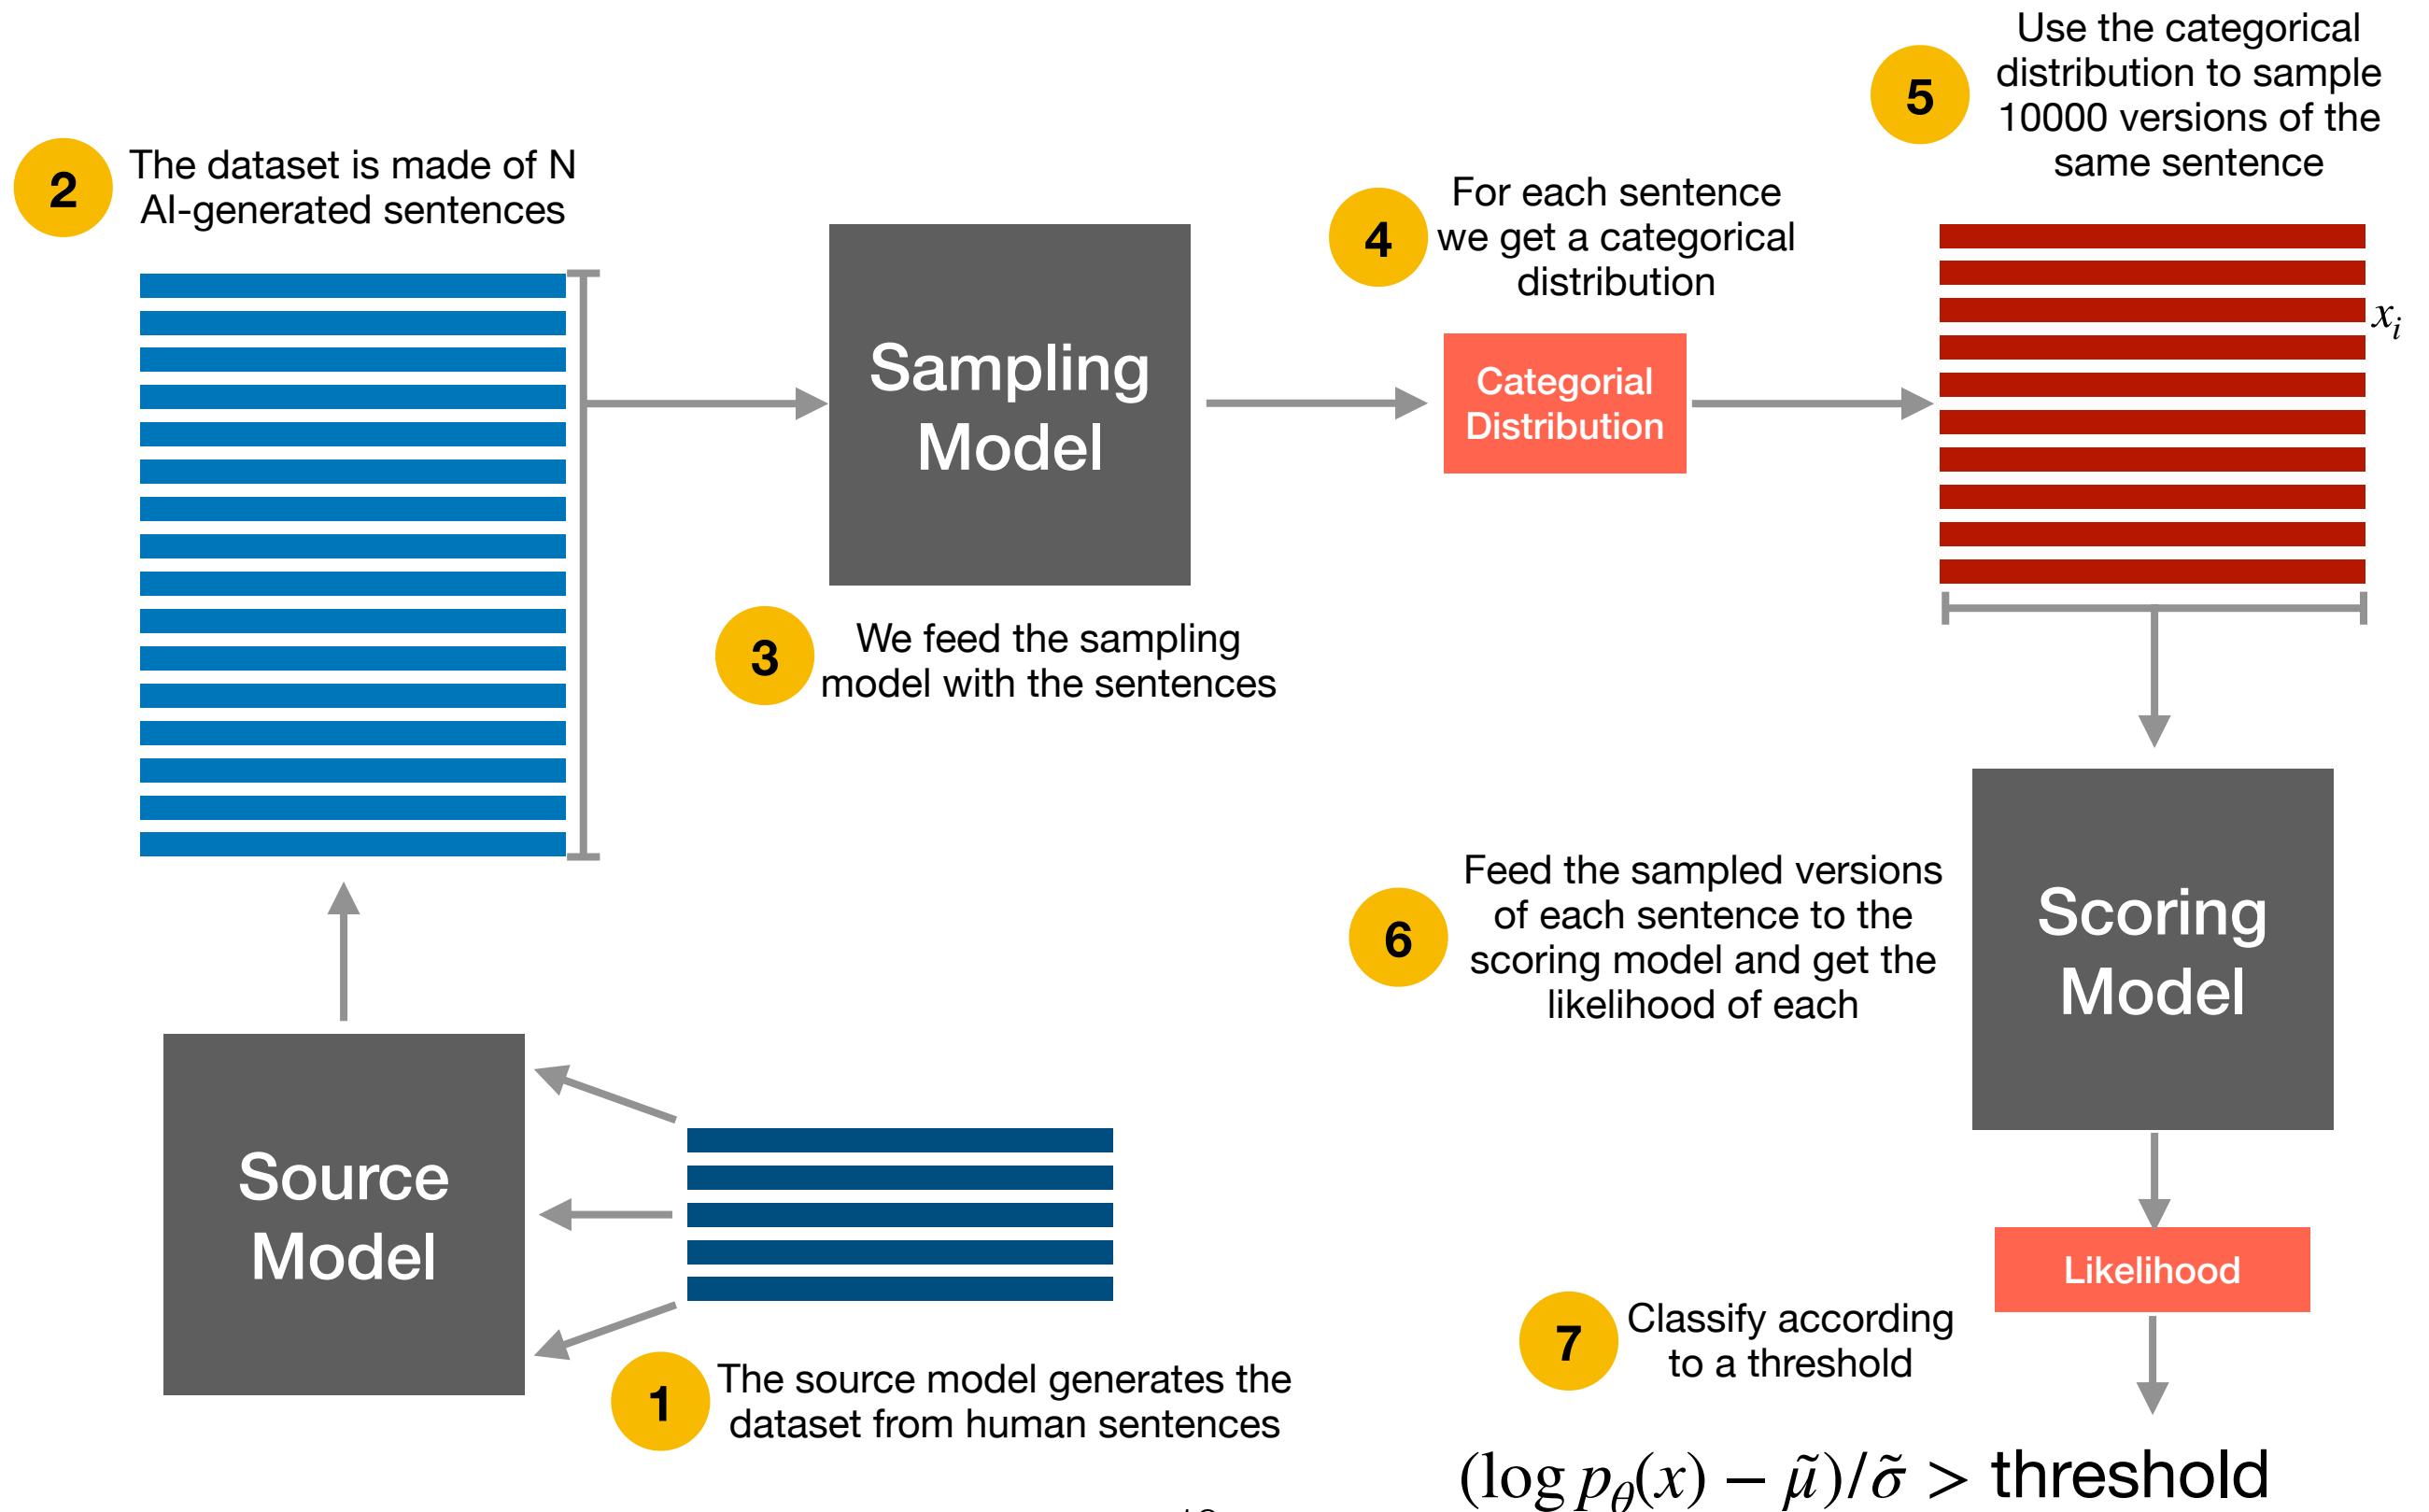
Baseline Methods




DetectGPT



FastDetectGPT



FastDetectGPT



1 AI generated sentence from the dataset originates **10000** samples and with those samples we compute this

FastDetectGPT

Why are we using this log-likelihood based sum?

$\log p_{\theta}(x_i)$

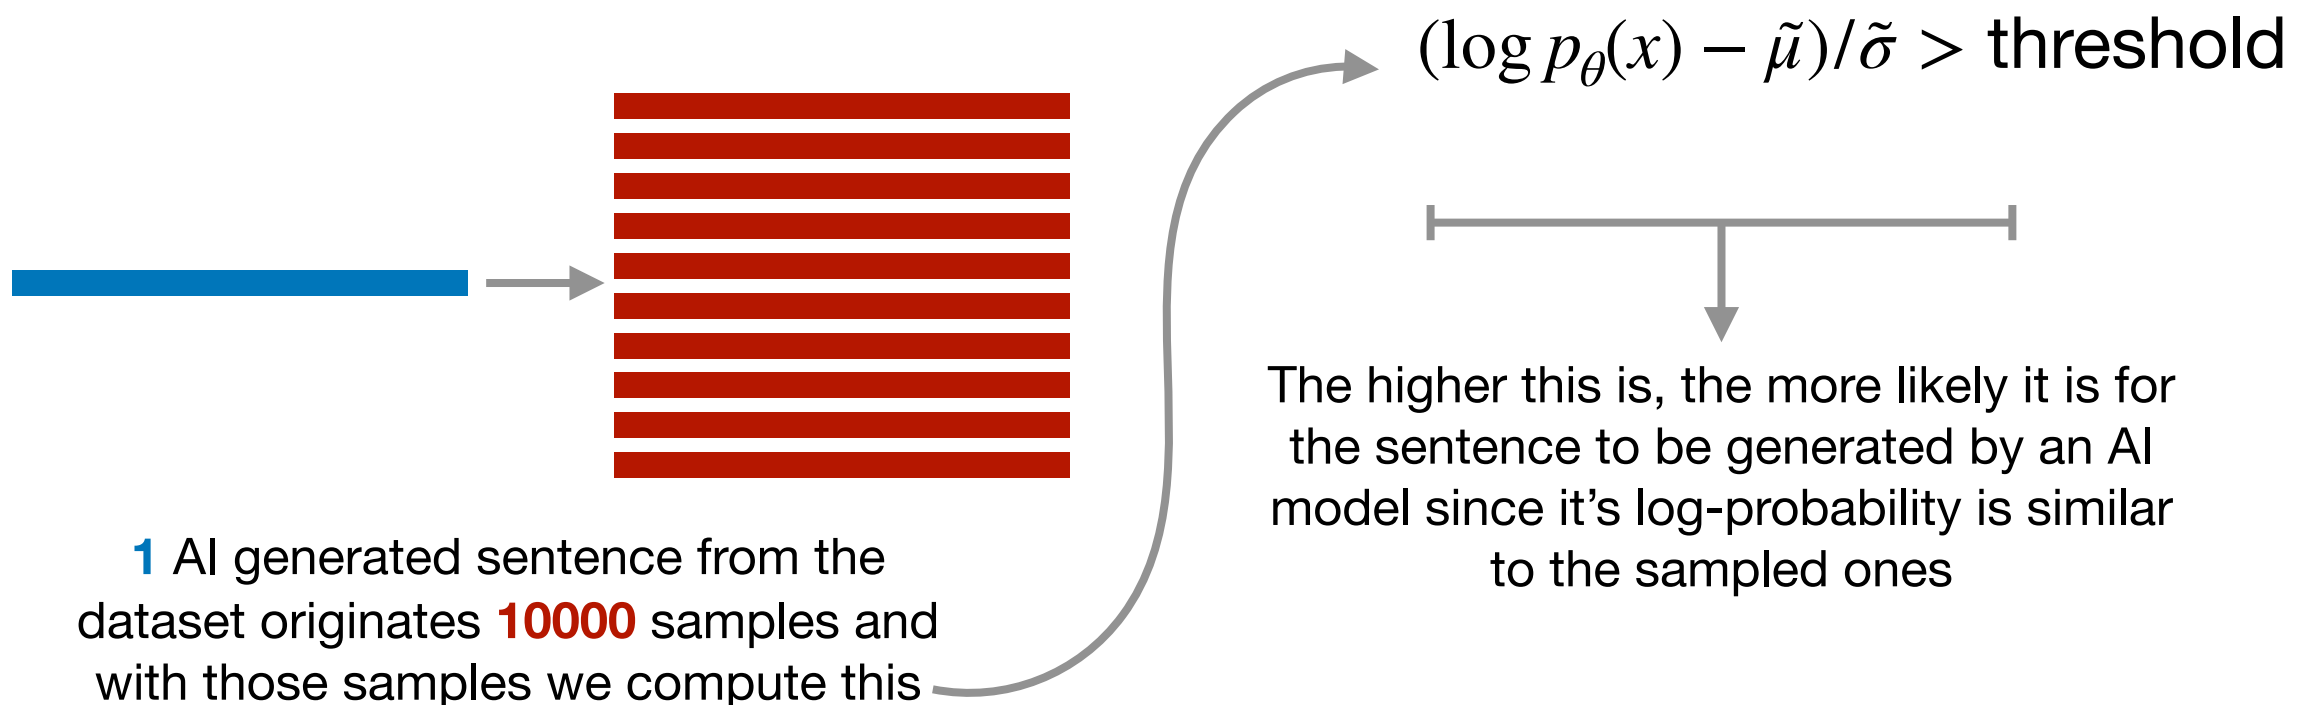
Is the log likelihood of the sentence x_i under the scoring model (sum of each token's log probability)

$\tilde{\mu}$

Is the mean of the log likelihood of each of the 10000 sample x_i

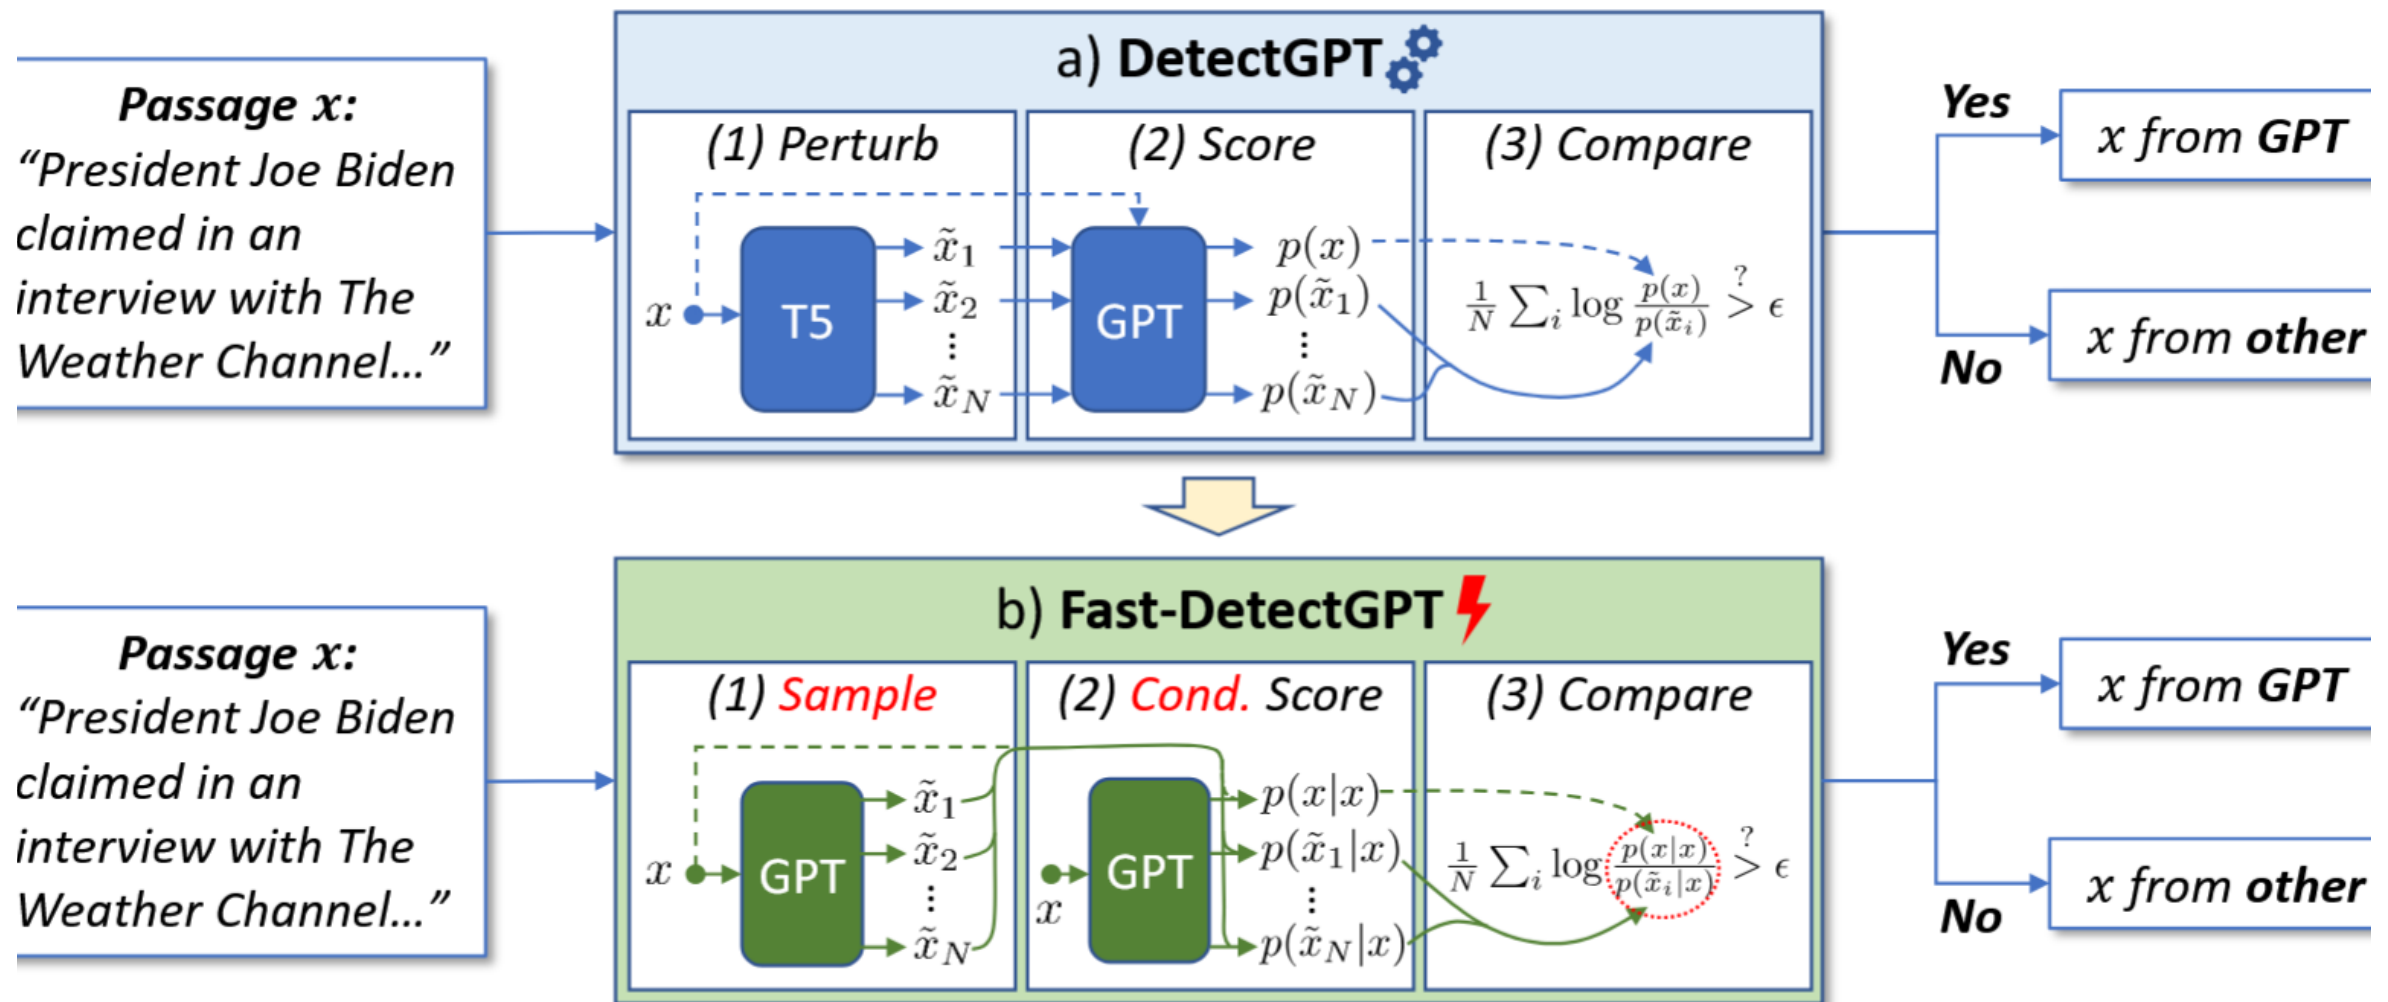
$\tilde{\sigma}$

Is the standard deviation of the log likelihood of each of the 10000 sample x_i



The Method

DetectGPT & Fast DetectGPT



Idea: perturb input text, feed to model & observe probability curvature

Baseline Models

Baselines = Threshold-based

```
# Baselines compute a single score
score = get_likelihood(model, text) # e.g., -2.5
score = get_rank(model, text)      # e.g., -150.2

# Then use a threshold for classification
if score > threshold:
    prediction = "AI-generated"
else:
    prediction = "Human-written"
```

DetectGPT/Fast-DetectGPT = Comparison-based

```
# DetectGPT compares original vs perturbations
original_ll = get_likelihood(model, original_text) # e.g., -2.5
perturbation_lls = [get_likelihood(model, pert) for pert in perturbations] # e.g.

# Z-score comparison (relative measure)
mean_pert = mean(perturbation_lls) # e.g., -2.93
std_pert = std(perturbation_lls)   # e.g., 0.15
z_score = (original_ll - mean_pert) / std_pert # (-2.5 - (-2.93)) / 0.15 = 2.87

# Higher z-score = more likely AI-generated
```

Key Difference:

Threshold approach:

- **Absolute judgment:** "This text has likelihood -2.5, is that high or low?"
- **Problem:** What's "high" depends on the model, domain, text length, etc.

Comparison approach:

FastDetectGPT

Here follows the pseudo-code for FastDetectGPT:

Algorithm 1 Fast-DetectGPT machine-generated text detection.

Input: passage x , sampling model q_φ , scoring model p_θ , and decision threshold ϵ

Output: True – probably machine-generated, False – probably human-written.


```
1: function FASTDETECTGPT( $x, q_\varphi, p_\theta$ )  
2:    $\tilde{x}_i \sim q_\varphi(\tilde{x}|x), i \in [1..N]$  ▷ Conditional sampling  
3:    $\tilde{\mu} \leftarrow \frac{1}{N} \sum_i \log p_\theta(\tilde{x}_i|x)$  ▷ Estimate the mean  
4:    $\tilde{\sigma}^2 \leftarrow \frac{1}{N-1} \sum_i (\log p_\theta(\tilde{x}_i|x) - \tilde{\mu})^2$  ▷ Estimate the variance  
5:    $\hat{d}_x \leftarrow (\log p_\theta(x) - \tilde{\mu})/\tilde{\sigma}$  ▷ Estimate conditional probability curvature  
6:   return  $\hat{d}_x > \epsilon$   
7: end function
```

$$\sum_i (\log p_\theta(x_i) - \tilde{\mu})/\tilde{\sigma} > \text{threshold}$$

Sampling and Scoring

Sampling and Scoring

Both sampling and scoring require the token distribution at each position in the sentence - sampling for alternative token selection, scoring for calculation of conditional probability curvature at each token



Generates alternative sequences by sampling from the model's token distribution



Scoring model used to compute conditional-probability curvature at each token; higher values mean the sequence is more likely machine-generated

Obtaining token distributions

We feed an LLM model the first m tokens of the original sentence in order to get logits for the next positions. After we have a probability distribution over each position on the sentence, we can build a categorical distribution with all the token distributions and use it either for sampling new sentences (sampling model) or scoring

Feed this into the model (first m tokens)

Let's

go

for

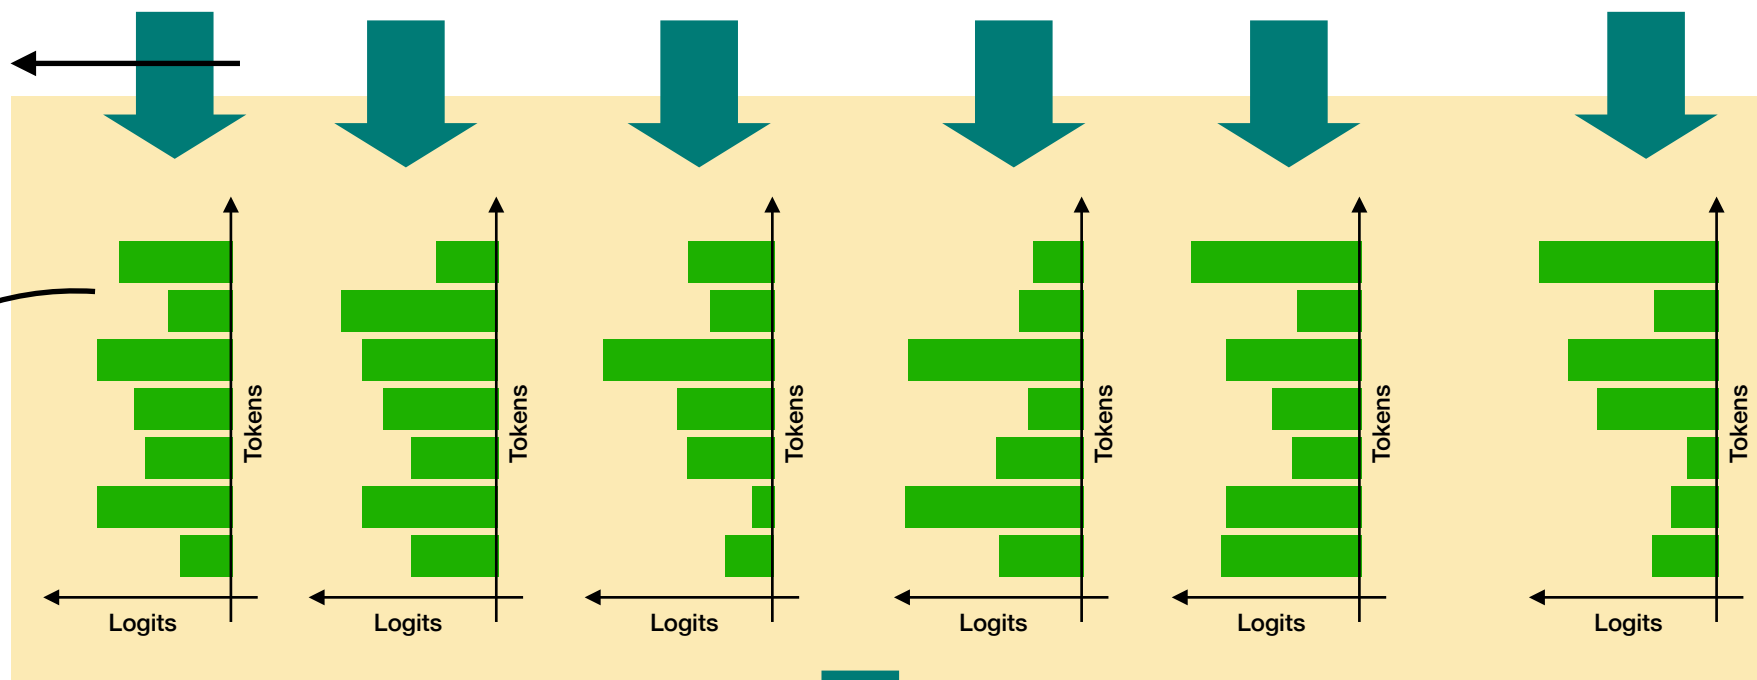
lunch

after

the

meeting

Get logits from white box LLM



Build a categorical distribution with all token distributions for the sampling model and for the scoring model

Obtaining token distributions

We feed an LLM model the first m tokens of the original sentence in order to get logits for the next positions. After we have a probability distribution over each position on the sentence, we can build a categorical distribution with all the token distributions and use it either for sampling new sentences (sampling model) or scoring

Feed this into the model (first m tokens)

Let's

go

for

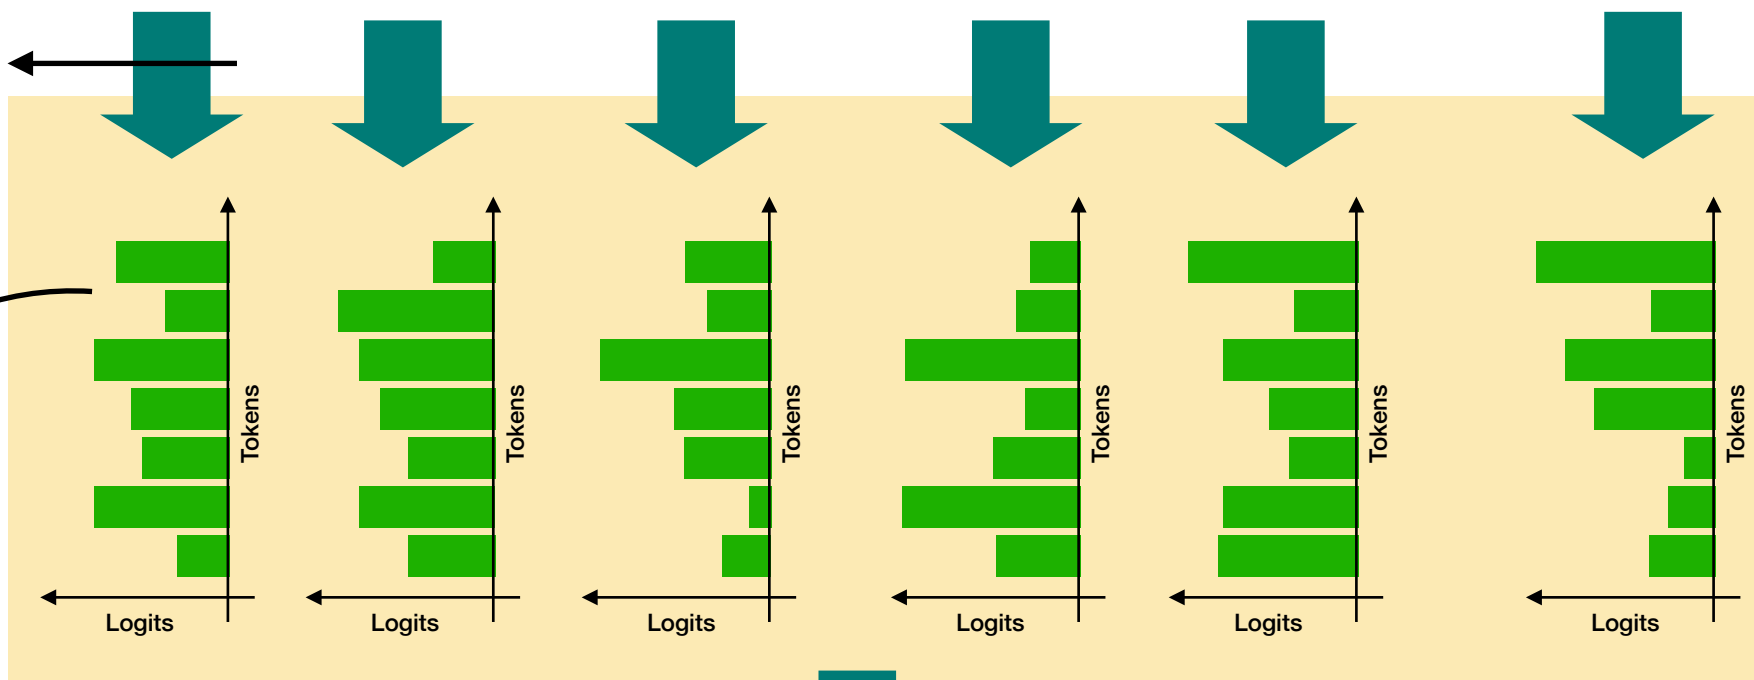
lunch

after

the

meeting

Get logits from white box LLM



Top-k Token distribution for this position given all previous tokens

Build a categorical distribution with all token distributions for the sampling model and for the scoring model

Each token's distributions is obtained based on the first m tokens of the original sentence

Probability of a generated sample

X

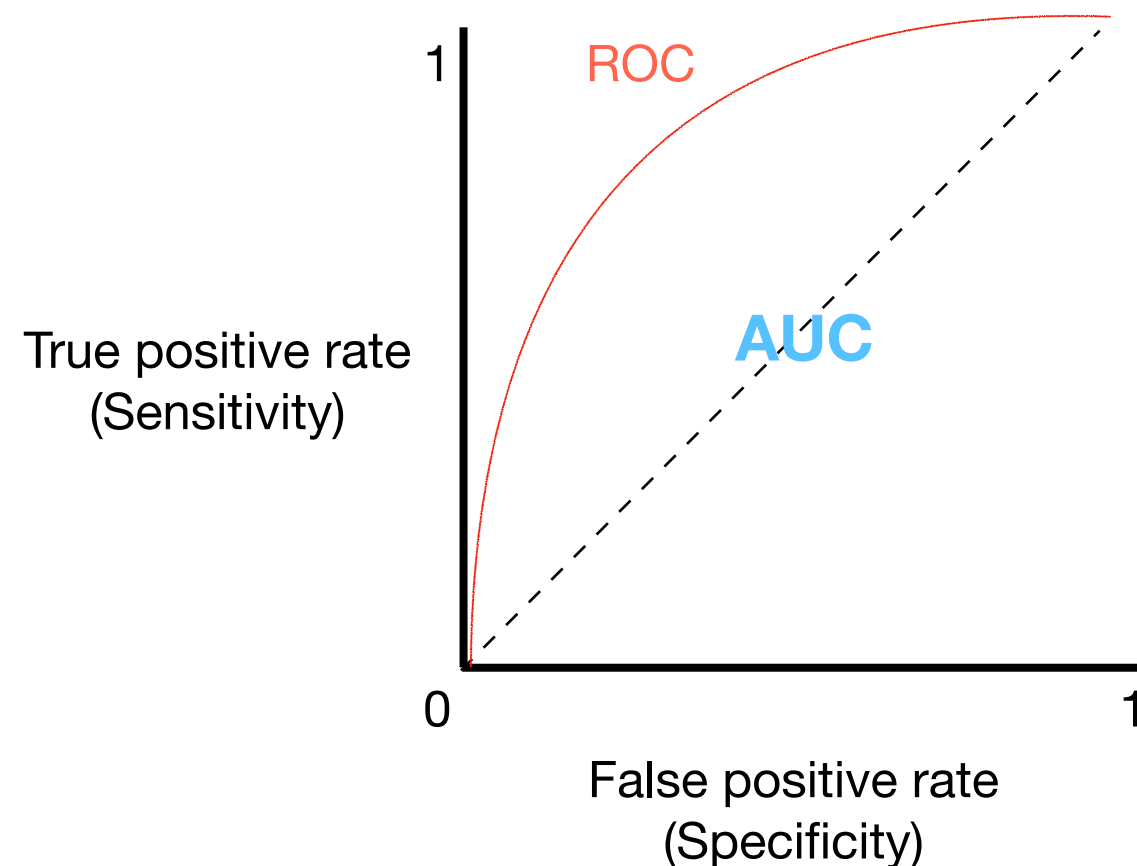
$$p_{\theta}(x^{\text{generated}} | x^{\text{original}}) = \prod_j p_{\theta}(x_j^{\text{generated}} | x_{<m}^{\text{original}})$$

The process of obtaining the tokens distribution for each position is done twice*: once for the **Sampling** Model and once for the **Scoring** Model

*Except for white box setting where both the models are **identical**.

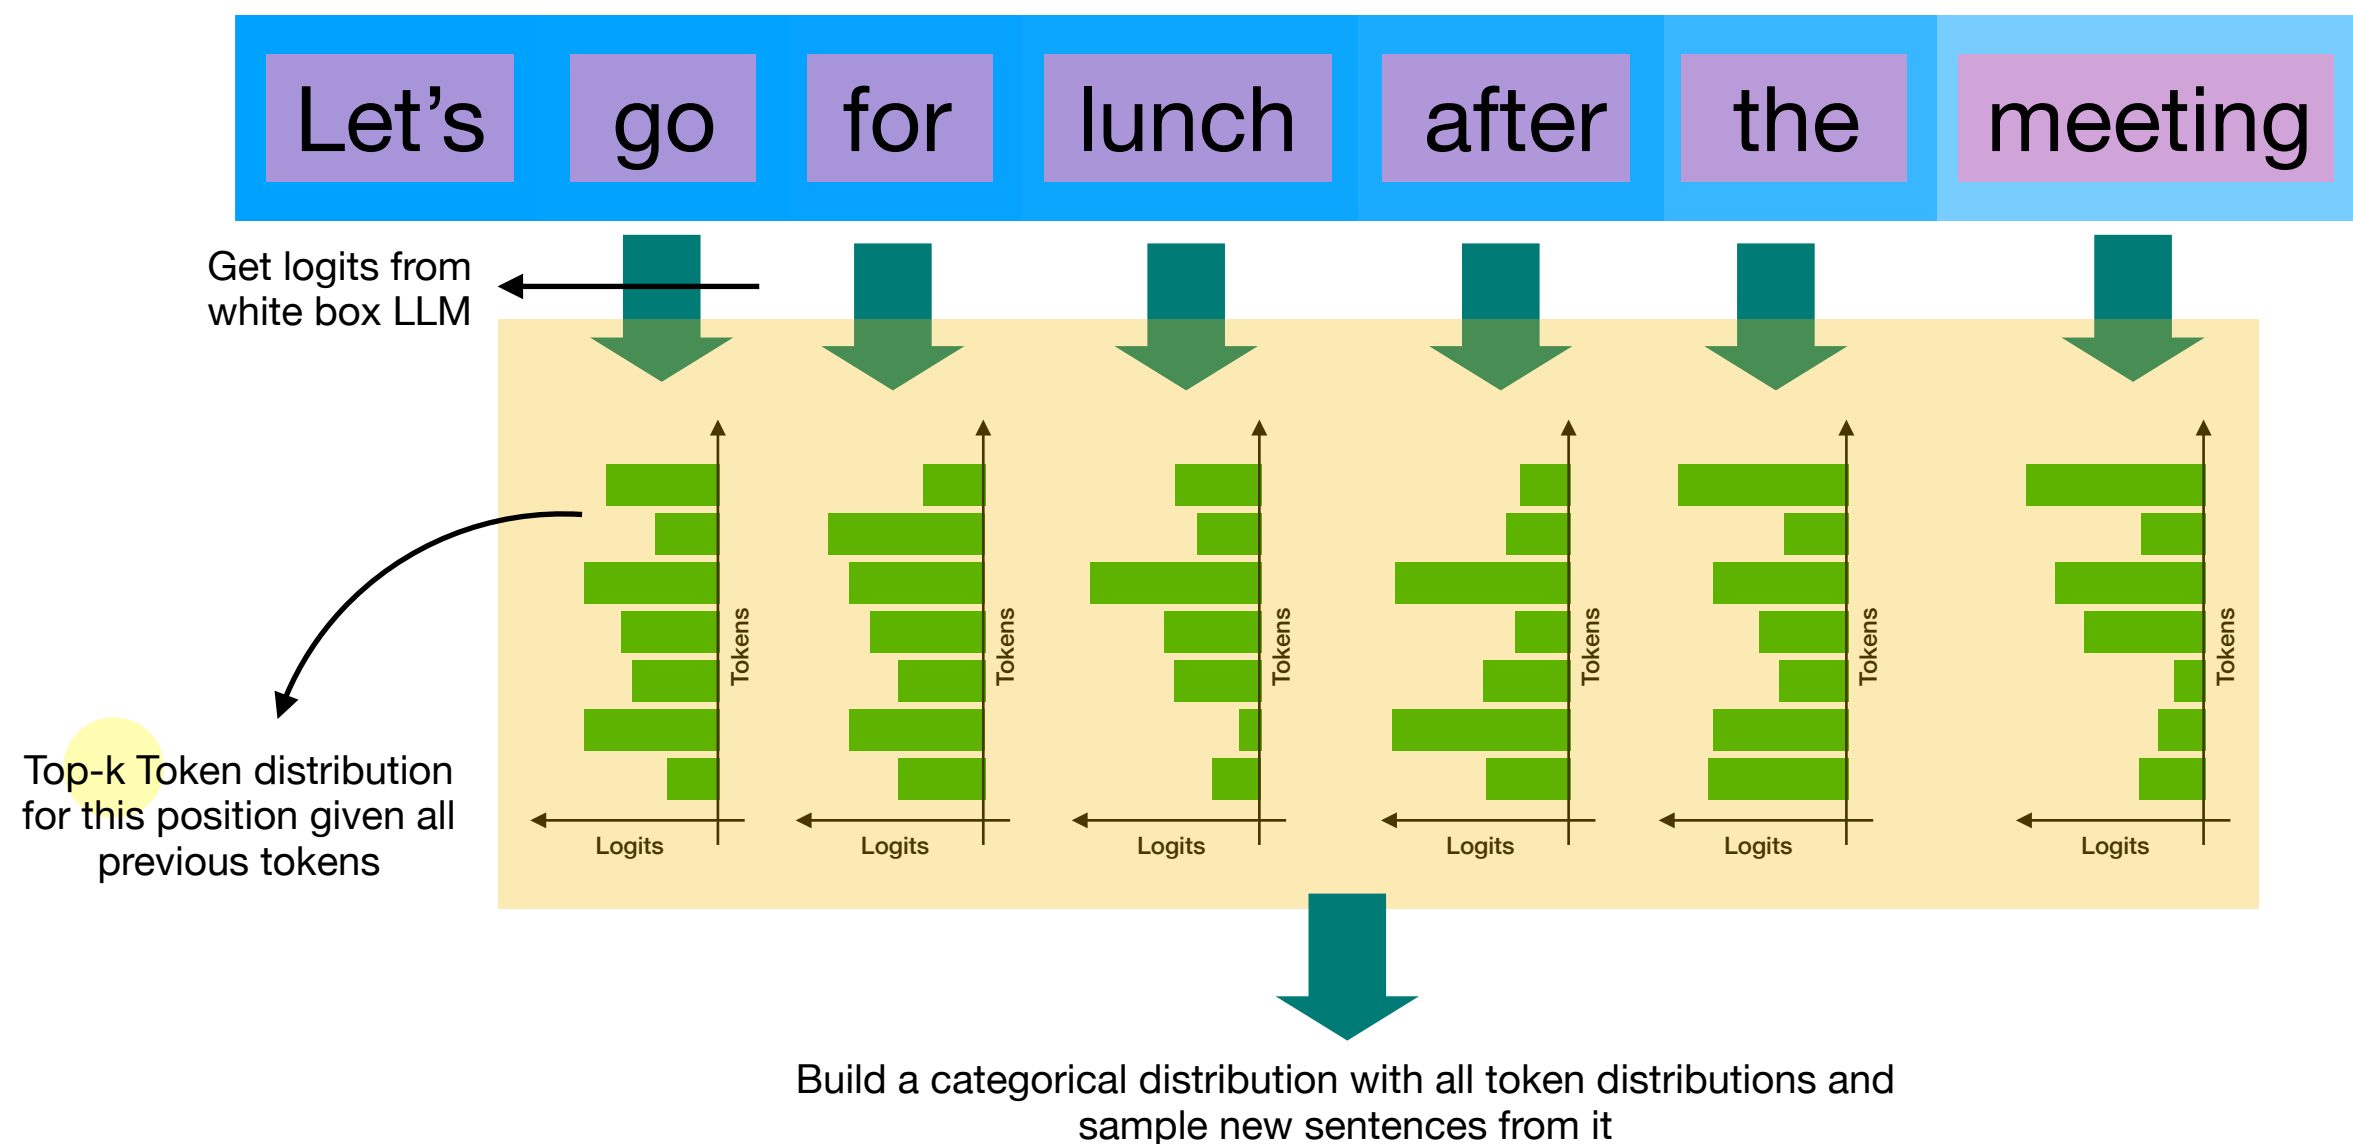
Main Metric

Area Under Receiver Operating Characteristic



How does sampling work? (Wrong)

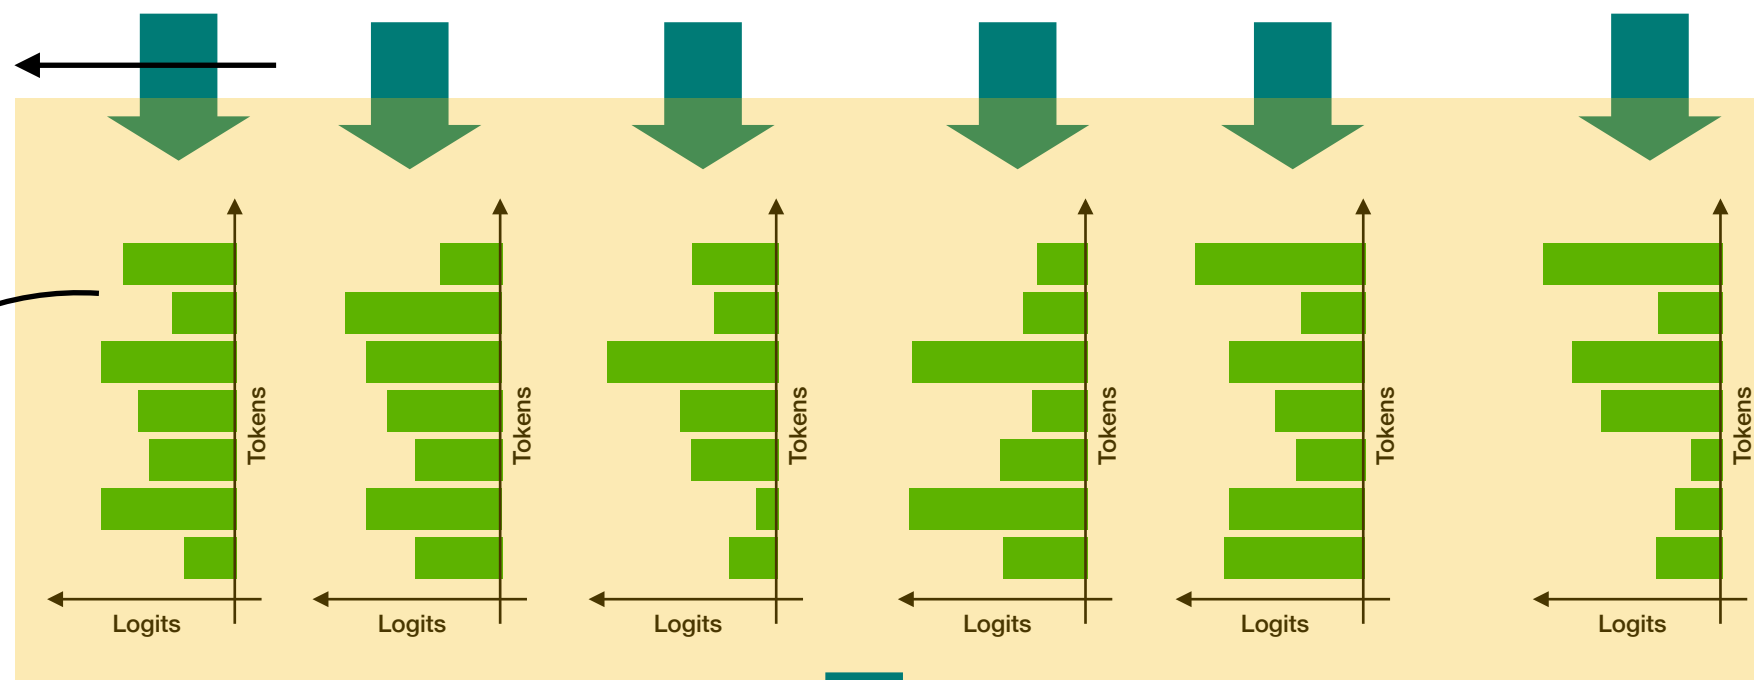
The goal is to build multiple versions of the same sentence (perturbed versions) and sampling is the way to do it. We feed an LLM model each sub-sentence of the original sentence in order to get logits for the next token. After we have a probability distribution over each token position on the sentence, we can build a categorical distribution with all the token distributions and sample from it in order to get a new sentence.



Let's go for lunch after the meeting

Get logits from white box LLM

Token distribution for this position given all previous tokens



Build a categorical distribution with all token distributions and sample new sentences from it

Let's go for lunch after the meeting

Experiments and Results

New Models and Datasets

Reproducibility Study

Small Open-Weight Models:

- DeepSeek R1 Distill Llama
- Microsoft Phi 2
- Mistral 7B Instruct v0.2

Metric: AUROC

Note: DetectGPT could only be used one time due to computational constraints

Method	XSum			
	R1-8B	Phi-2	Mistral	Avg.
Likelihood	0.9999	0.8782	0.9665	0.9482
Entropy	0.1645	0.5411	0.4558	0.3871
LogRank	1.0000	0.9051	0.9596	0.9549
DetectGPT	0.8560	-	-	-
★ Fast-DetectGPT	1.0000	0.9769	0.9989	0.9919

Method	HC3			
	R1-8B	Phi-2	Mistral	Avg.
Likelihood	0.8391	0.7139	0.7486	0.7672
Entropy	0.3521	0.4358	0.4235	0.4038
LogRank	0.8275	0.7224	0.7393	0.7631
★ Fast-DetectGPT	0.9710	0.8347	0.8963	0.9007

Speed Comparison

Fast-DetectGPT is

118.9x

faster than DetectGPT

on XSum with DeepSeek R1 Distill Llama

Close-Weight API Model

Method	GPT 4o mini		
	XSum ★	HC3	Avg.
Fast-DetectGPT	0.8051	0.6026	0.7039

- Lower performance due to top-20 probability restriction

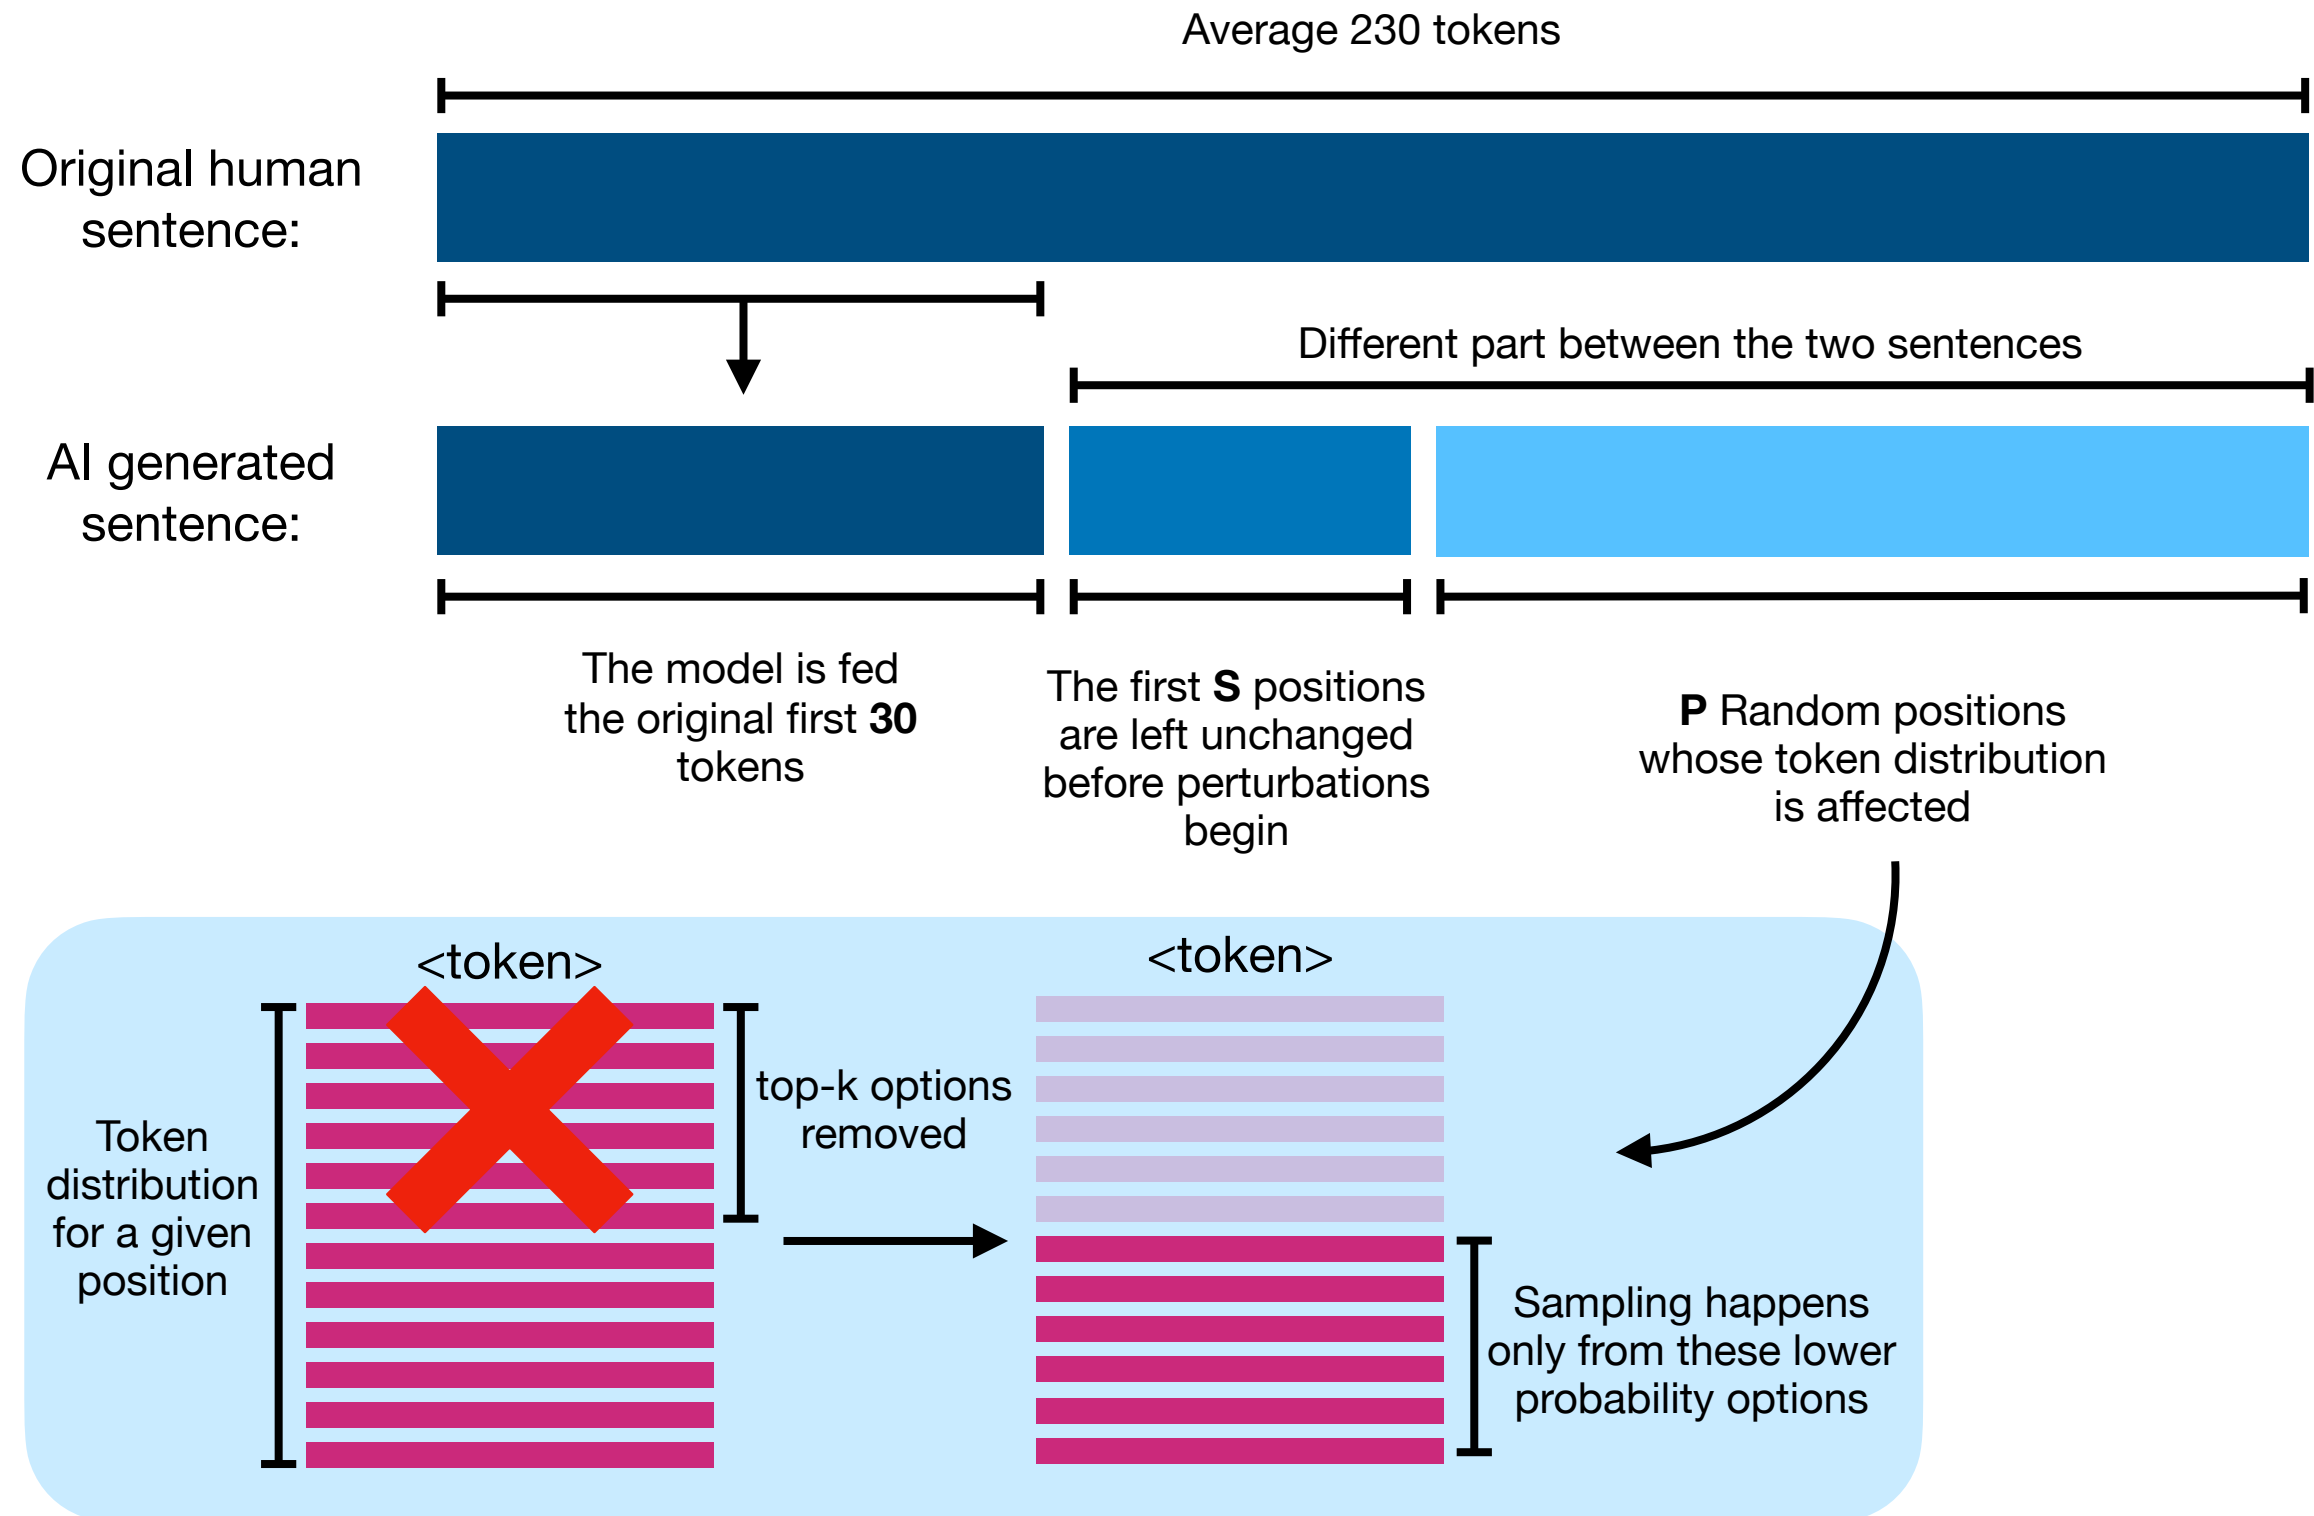
Low-Resource Language

Method	Underrepresented language			
	R1-8B	Phi-2	Mistral	Avg.
Likelihood	1.0000	0.9976	0.9972	0.9983
Entropy	0.0010	0.0968	0.2730	0.1236
★ LogRank	1.0000	0.9984	0.9974	0.9986
Fast-DetectGPT	0.9994	0.9825	0.9995	0.9938


Tested on Scottish Gaelic subset from XL-Sum
+ Competitive scores, overall robustness

Mid-K Extension

How Mid-K Works













Mid-K in Action

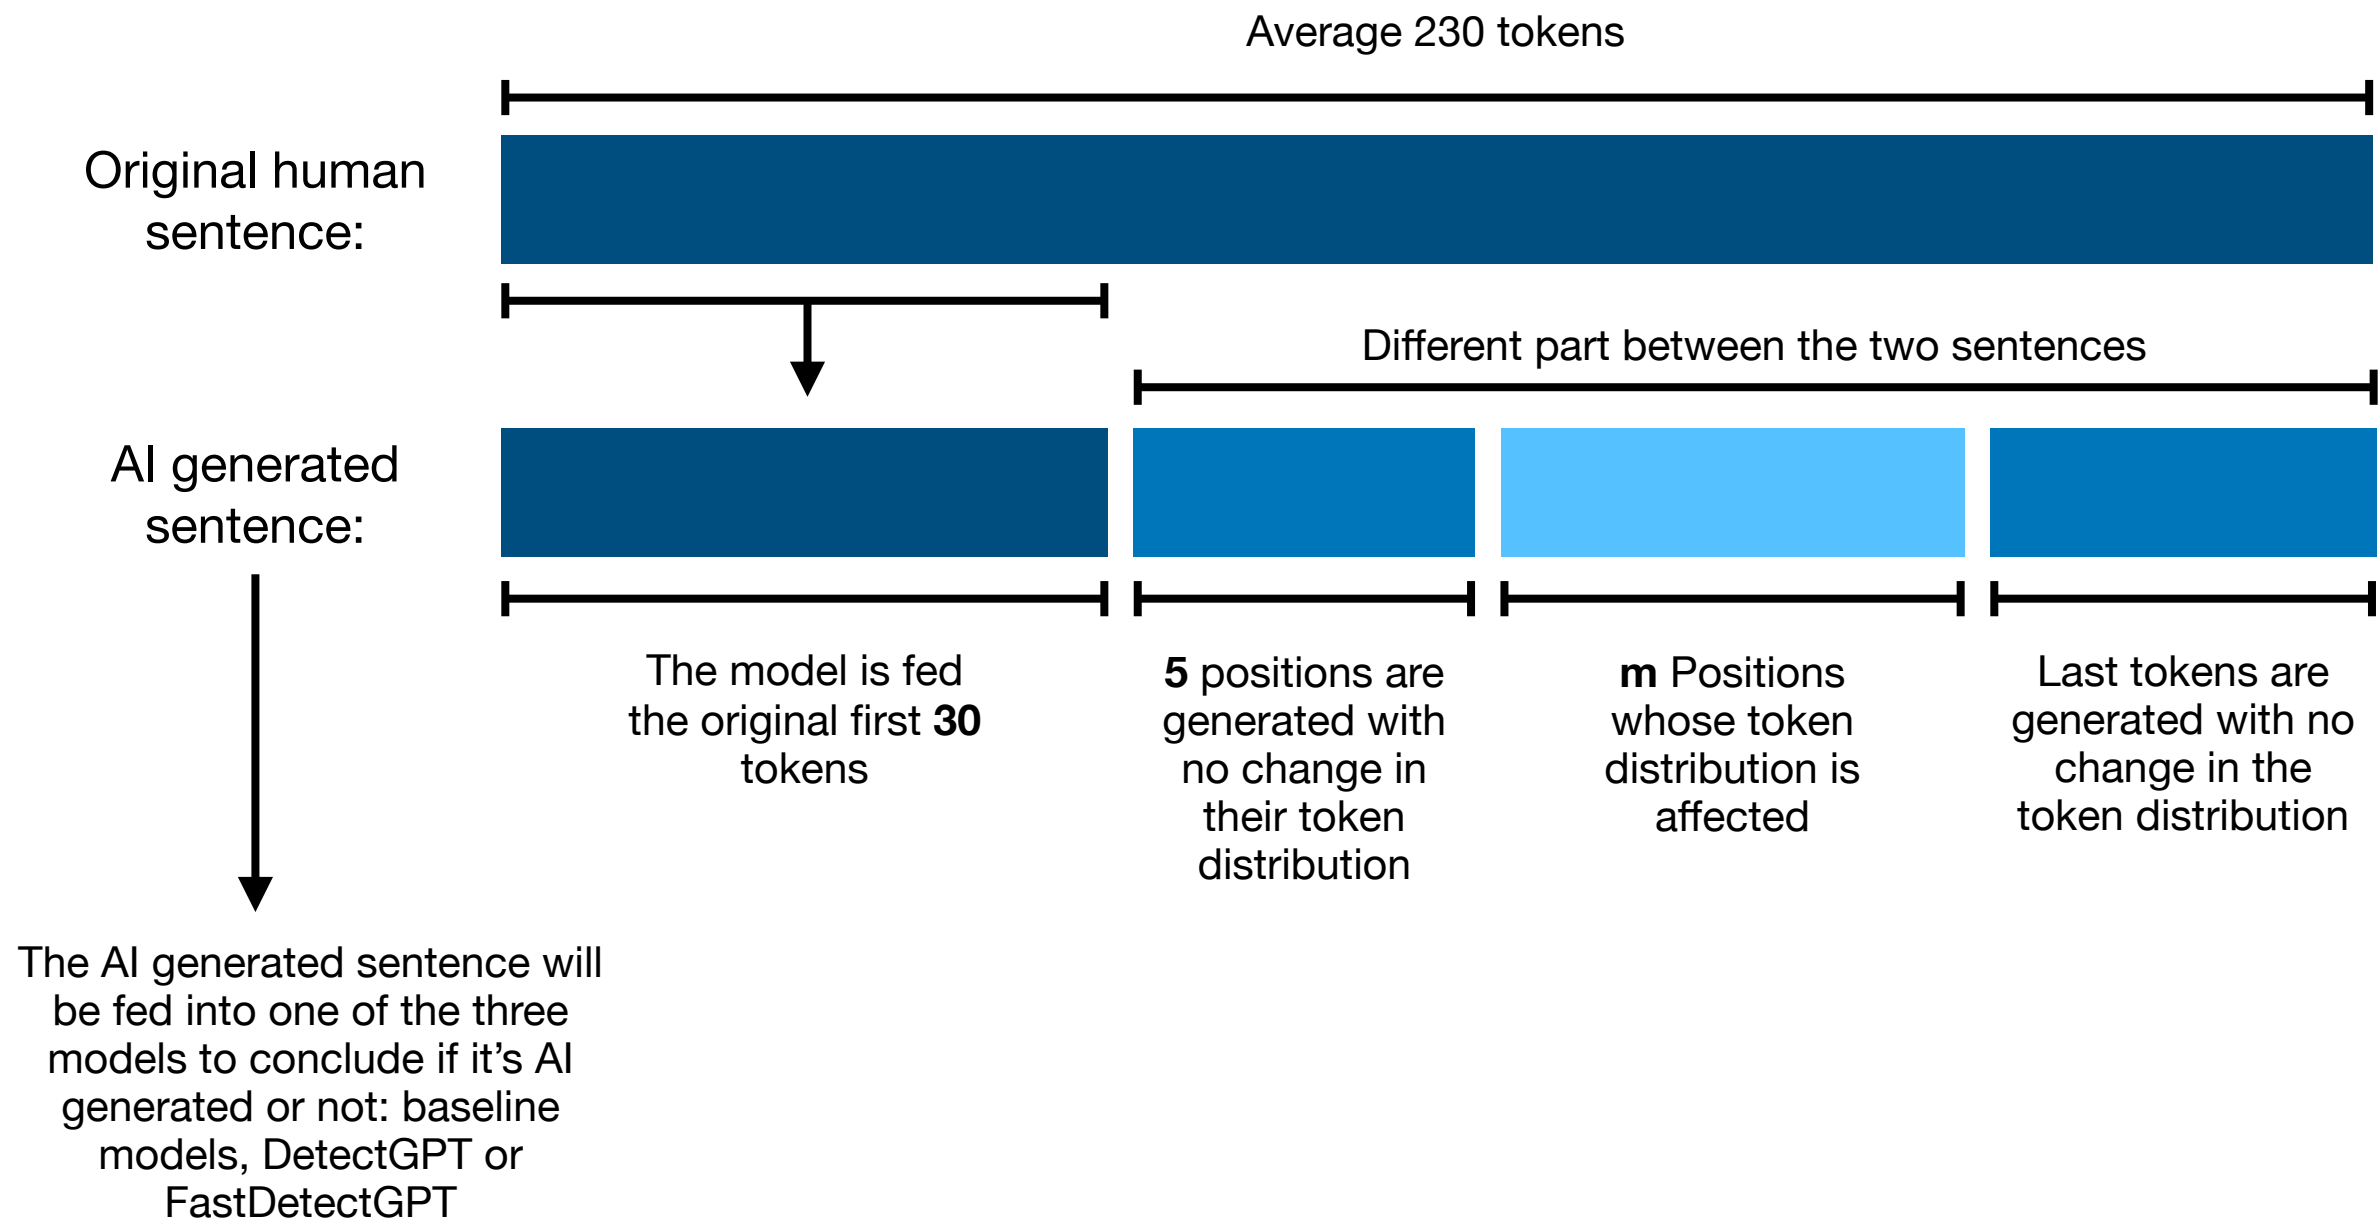


- Key parameters:
 - **P**: Number of random positions where attack is applied
 - **K**: Number of top tokens removed at each position
 - **S**: Start position (attack begins after S tokens)



Examples

Position	1	...	30	31	...	35	36	~250
Token	<token>		<token>	<token>	<token>	<token>	<token>	<token>	<token>	<token>	<token>	<token>	<token>
State	Original	Original	Original	Sampled	Sampled	Sampled	Sampled	Sampled	Sampled	Sampled	Sampled	Sampled	Sampled
													

Mid-k dataset generation



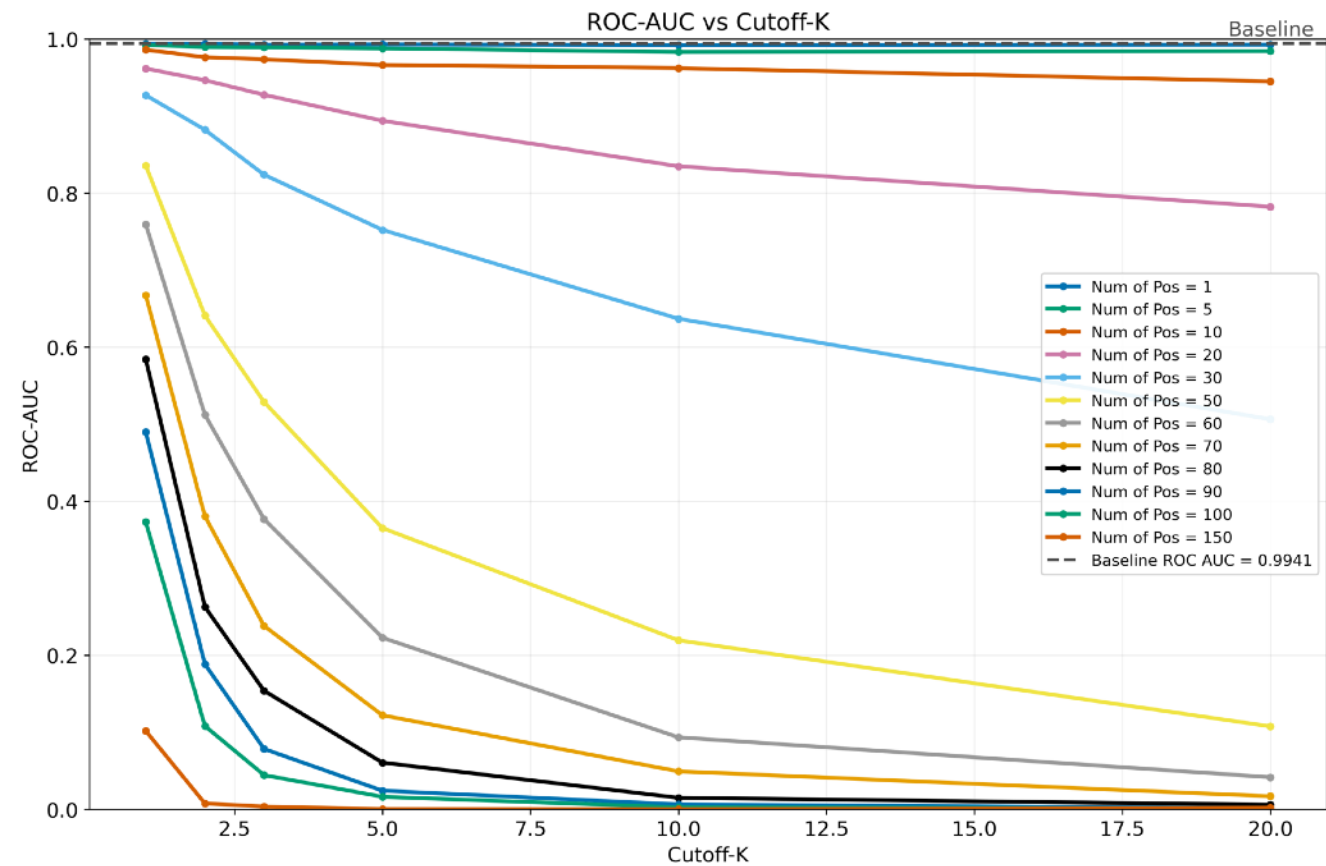
Effect of Mid-K Attack



Setup:

- Model: GPT-2
- Dataset: XSum
- Parameters: $P \in \{1, 5, 10, 20, 30, 50, 60, 70, 80, 90, 100, 150\}$
- $K \in \{1, 2, 5, 10, 20\}$, $S = 5$

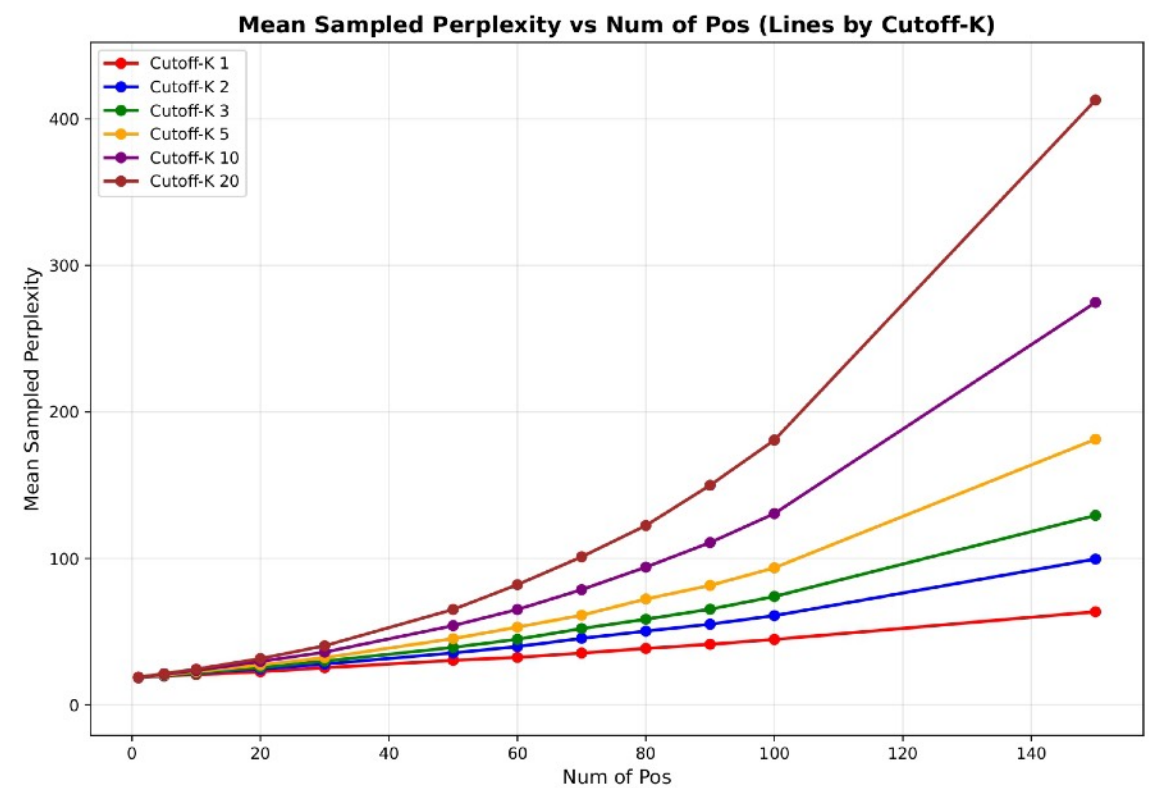

Effect of Mid-K Attack



Setup:

- Model: GPT-2
- Dataset: XSum
- Parameters: $P \in \{1, 5, 10, 20, 30, 50, 60, 70, 80, 90, 100, 150\}$
- $K \in \{1, 2, 5, 10, 20\}$, $S = 5$

Effect of Mid-K Attack



Using a dataset where the AI generated sentences have tokens with lower log probability brings the whole log-likelihood of the generated sentence down. This will trick FastDetectGPT because it focuses on concluding that sentences with high log-probability are AI-generated. In this case, they are still AI-generated but the log-probability is lower, that's why the AUROC goes down as we increase the number of positions with removed top-m logits and as we increase m.

Qualitative Examples

Fixed part: "The false message said: "Breaking: Two Explosions..."

Human	in the White House and Barack Obama is injured." US markets were spooked by the tweet
AI (No Attack)	At University of Michigan Student Residences, Campus Police Service (APS) has confirmed at least one person was hurt.
Mild (P=60, K=2)	At University of Michigan Student Residences, Campus Police Service (APS) has confirmed at 7:18 a.b.'s report a report by Campus Police that three occupants were shot dead
Wide (P=150, K=1)	At University of Michigan Student Residences, Campus Fires . We apologize if we didn't find the wrong tweet. The University has since been informed.
Strong (P=150, K=20)	are Found at a School; Parents Support Public Alert: Close Education Week Hours Soon With Two Public Works, Infrastructure Flashed". http@worldc.en

Perplexity

AUROC		Number of positions with top-m logits being removed			
		5	10	20	30
Number of logits removed from the probability distribution (m)	1				
	3				
	5				
	10				
	20				

Conclusion

Takeaways

Fast-DetectGPT:

- Outperforms baselines on newly tested models ✓
- Two orders of magnitude faster than DetectGPT (not 340 times though) ✓
- Vulnerable towards adversarial attacks ✗
- Performance depends on full logits distribution availability ✗

Thank you!